

0.13.0

Table of contents

Mission Modeler Guide

[Foundations of Simulation & Modeling](#)

[Simulation Mechanics](#)

[Creating a Mission Simulation Overview](#)

[Developing a Mission Model](#)

[Configuring a Mission Model](#)

[Activities](#)

[Activity Mappers](#)

[Value Schemas](#)

[Declaring Parameters](#)

[Models & Resources](#)

[Creating Plans](#)

[Constraints](#)

[Precomputed Resource Profiles](#)

[Scheduling](#)

Command Expansion Guide

[Upload and Generate Command Typescript Library](#)

[Generate Activity Typescript Library](#)

[IDE Setup](#)

[Expansion Logic API](#)

[Submit Expansion Logic](#)

[Create an Expansion Set](#)

[Expand the Plan](#)

[Get Expanded Commands](#)

[Create Sequence](#)

[Link Simulated Activity to Sequence](#)

[Retrieve Sequence in SeqJson](#)

[Generate a standalone sequence in SeqJson](#)

[Valid Sequence JSON Format](#)

User Guide

[Merlin Activity Plans](#)

[Aerie GraphQL API](#)

[Simulation Results](#)

[Planning UI](#)

[Simulation Configuration](#)

[UI Views](#)

[Activity Directive Metadata](#)

Product Guide

[Installation](#)

[System Requirements](#)

[Administration](#)

[Support](#)

Mission Modeler Guide

Introduction

This guide explains how to make use of Aerie-provided capabilities in the latest version of Aerie, and will be updated as Aerie evolves.

Aerie is a new software system being developed by the MPSA element of MGSS (Multi-mission Ground System and Services), a subsystem of AMMOS (Advanced Multi-mission Operations System). Aerie will support mission operations by providing capabilities for activity planning, sequencing, and spacecraft analysis. These capabilities include modeling & simulation, scheduling, and validation. Aerie will replace several legacy MGSS tools, including but not limited to APGEN, SEQGEN, MPS Editor and MPS Server.

Aerie currently provides the following elements.

- Merlin modeling framework for defining mission resources and activity types.
- Merlin web UI for activity planning and simulation analysis.
- Falcon sequence editor UI

Mission Modeling with the Merlin Framework

In Merlin, a mission model serves activity planning needs in two ways. First, it describes how various mission resources behave autonomously over time. Second, it defines how activities perturb these resources at discrete time points, causing them to change their behavior. This information enables Aerie to provide scheduling, constraint validation, and resource plotting capabilities on top of a mission model.

The Merlin Framework empowers adaptation engineers to serve the needs of mission planners maintain as well as keep their codebase maintainable and testable over the span of a mission. The Framework aims to make the experience of mission modeling similar to standard Java development, while still addressing the unique needs of the simulation domain.

In the Merlin Framework, a mission model breaks down into two types of entity: system models and activity types.

System models can range in complexity from a single aspect of an instrument to an entire mission. In fact, Merlin only requires one system model to exist: the top-level mission model. The mission model can delegate to other, more focused models, such as subsystem models, which may themselves delegate further. Ultimately, fine-grained models capture the system state they own in a `Cell`, which is a simulation-aware analogue of a Java mutable field. (**In Merlin, mutable fields on models must not be used.** All mutable state must be controlled by a `Cell`.) Models may provide regular Java methods for interacting with that state, and other models (including activity types) may invoke those methods.

Activity types are a specialized kind of model. Each activity type defines the parameters for activities of that type, which may be instantiated and configured by a mission planner. An activity type also defines a single method that acts as the entrypoint into the simulated system: when an activity of that type occurs, its method is invoked with the activity parameters and the top-level mission model. It may then interact freely with the rest of the system.

Just as activity types define the entrypoints into a simulation, the mission model also defines *resources*, which allow information to be extracted from the simulation. A resource is associated with a method that returns a "dynamics" -- a description of the current autonomous behavior of the resource. Merlin currently provides discrete dynamics (constants held over time) and linear dynamics (real values varying linearly with time), and is designed to support more in the future.

A simulation over a mission model iteratively runs activities and queries resources for their updated dynamics, producing a composite profile of dynamics for each resource over the entire simulation duration.

Simulation Mechanics

In progress...

Creating a Mission Simulation Overview

Running an Aerie simulation requires a mission model that describes effects of activities over modeled resources, and a plan file that declares a schedule of activity instances with specified parameters. A plan file must be created with respect to an mission model file, since activities in a plan and their parameters are validated against activity type definitions in the mission model.

Here's a summary workflow to getting a simulation result from Aerie.

1. Install Aerie services following instructions on [Product Guide](#)
2. Create a Merlin mission model following the instructions on [Developing a Mission Model](#) page.
3. Upload the mission model to Aerie through [Planning Web GUI](#) or [Aerie API GraphQL graphical interface](#).
4. Create a plan associated with the mission model using again the [Planning Web GUI](#) or [Aerie API](#)
5. Trigger simulation via either interfaces listed above.

Developing a Mission Model

A mission model defines the behavior of any measurable mission resources, and a set of **activity types**, defining the ways in which a plan may influence mission resources.

Mission Modeling Libraries

Aerie provides seven libraries that can be imported by a project. They are:

- [contrib](#)
- [merlin-framework](#)
- [merlin-framework-processor](#)
- [merlin-framework-junit](#)
- [merlin-sdk](#)
- [merlin-driver](#)
- [parsing-utilities](#)

Package-info File

A mission model must contain, at the very least, a `package-info.java` containing annotations that describe the highest-level features of the mission model. For example:

```
// examples/banananation/package-info.java
@MissionModel(model = Mission.class)
@WithActivityType(BiteBananaActivity.class)
@WithActivityType(PeelBananaActivity.class)
@WithActivityType(ParameterTestActivity.class)
@WithMappers(BasicValueMappers.class)
package gov.nasa.jpl.aerie.banananation;

import gov.nasa.jpl.aerie.banananation.activities.BiteBananaActivity;
import gov.nasa.jpl.aerie.banananation.activities.ParameterTestActivity;
import gov.nasa.jpl.aerie.banananation.activities.PeelBananaActivity;
import gov.nasa.jpl.aerie.contrib.serialization.rulesets.BasicValueMappers;
import gov.nasa.jpl.aerie.merlin.framework.annotations.MissionModel;
import gov.nasa.jpl.aerie.merlin.framework.annotations.MissionModel.WithActivityType;
import gov.nasa.jpl.aerie.merlin.framework.annotations.MissionModel.WithMappers;
```

This `package-info.java` identifies the top-level class representing the mission model, and registers activity types that may interact with the mission model. Merlin processes these annotations at compile-time, generating a set of boilerplate classes which take care of interacting with the Aerie platform.

The `@WithMappers` annotation informs the annotation processor of a set of serialization rules for activity parameters of various types; the `BasicValueMappers` ruleset covers most primitive Java types. Mission modelers may also create their own rulesets, specifying rules for mapping custom value types. If multiple mapper classes are included via the `@WithMappers` annotations, and multiple mappers declare a mapping rule to the same data type, the rule found in the earlier declared mapper will take precedence. For more information on allowing custom values, see [value mappers](#).

Mission Model Class

The top-level mission model is responsible for defining all of the mission resources and their behavior when affected by activities. Of course, the top-level model may delegate to smaller, more focused models based on the needs of the mission. The top-level model is received by activities, however, so it must make accessible any resources or methods to be used therein.

```
// examples/banananation/Mission.java
public class Mission {
    public final AdditiveRegister fruit = AdditiveRegister.create(4.0);
    public final AdditiveRegister peel = AdditiveRegister.create(4.0);
    public final Register<Flag> flag = Register.create(Flag.A);

    public Mission(final Registrar registrar, final Instant planStart) {
        registrar.discrete("/flag", this.flag, new EnumValueMapper<>(Flag.class));
        registrar.real("/peel", this.peel);
        registrar.real("/fruit", this.fruit);
    }
}
```

If the second constructor parameter is an `Instant` then the plan start time will be passed to this constructor. This is an entirely optional parameter. For example, the following top-level model is also perfectly acceptable:

```
// examples/banananation/Mission.java
public class Mission {
    public final AdditiveRegister fruit = AdditiveRegister.create(4.0);
    public final AdditiveRegister peel = AdditiveRegister.create(4.0);
    public final Register<Flag> flag = Register.create(Flag.A);

    public Mission(final Registrar registrar) {
        registrar.discrete("/flag", this.flag, new EnumValueMapper<>(Flag.class));
        registrar.real("/peel", this.peel);
        registrar.real("/fruit", this.fruit);
    }
}
```

Mission resources are declared using `Registrar#discrete()` or `Registrar#real()`.

A model may also express autonomous behaviors, where a discrete change occurs in the system outside of an activity's effects. A **daemon task** can be used to model these behaviors. Daemons are spawned at the beginning of any simulation, and may perform the same effects as an activity. Daemons are prepared using the `spawn()` method.

Activity types

An **activity type** defines a simulated behavior that may be invoked by a planner, separate from the autonomous behavior of the mission model itself. Activity types may define **parameters**, which are filled with **arguments** by a planner and provided to the activity upon execution. Activity types may also define **validations** for the purpose of informing a planner when the parameters they have provided may be problematic.

```
// examples/banananation/activities/PeelBananaActivity.java
@ActivityType("PeelBanana")
public final class PeelBananaActivity {
    private static final double MASHED_BANANA_AMOUNT = 1.0;

    @Parameter
    public String peelDirection = "fromStem";

    @Validation("peel direction must be fromStem or fromTip")
    public boolean validatePeelDirection() {
        return List.of("fromStem", "fromTip").contains(this.peelDirection);
    }

    @EffectModel
    public void run(final Mission mission) {
        if (peelDirection.equals("fromStem")) {
            mission.fruit.subtract(MASHED_BANANA_AMOUNT);
        }
        mission.peel.subtract(1.0);
    }
}
```

Merlin automatically generates parameter serialization boilerplate for every activity type defined in the mission model's `package-info.java`. Moreover, the generated `Model` base class provides helper methods for spawning each type of activity as children from other activities.

Uploading a Mission Model

In order to use a mission model to simulate a plan on the Aerie platform, it must be packaged as a JAR file with all of its non-Merlin dependencies bundled in. The [template mission model](#) provides this capability out of the box, so long as your dependencies are specified with Gradle's `implementation` dependency class. The built mission model JAR can be uploaded to Aerie through the Aerie web UI.

Configuring a Mission Model

A **mission model configuration** enables mission modelers to set initial mission model values when running a simulation. Configurations are tied to a plan, therefore each plan is able to define its own set of configuration parameters. The `examples/banananation` project contains a `Configuration` data class example to demonstrate how a simple configuration may be created.

Setup

Mission Model

The `examples/banananation` project makes use of the `@WithConfiguration` annotation within `package-info.java`:

```
@MissionModel(model = Mission.class)

@WithConfiguration(Configuration.class)
```

When the `@WithConfiguration` annotation is used, the model – defined within the `@MissionModel` annotation – must accept the configuration as the last constructor parameter. See `Mission.java`:

```
public Mission(final Registrar registrar, final Configuration config) {
    // ...
}
```

If the second argument is an `Instant` denoting plan start time the configuration parameter must still be provided as the last argument. For example:

```
public Mission(final Registrar registrar, final Instant planStart, final Configuration config) {
    // ...
}
```

Configuration

A configuration class should be defined with the same parameter annotations as activities. See [Declaring Parameters](#) for a thorough explanation of all possible styles of `@Export` parameter declaration and validation.

Similarly to activities, the Merlin annotation processor will take care of all serialization/deserialization of the configuration object. The Merlin annotation processor will generate a configuration mapper for the configuration defined within `@WithConfiguration()`.

Examples

`Configuration` can be a simple data class. For example:

```
package gov.nasa.jpl.aerie.banananation;

import java.nio.file.Path;

import static gov.nasa.jpl.aerie.merlin.framework.annotations.Export.Template;

public record Configuration(int initialPlantCount, String initialProducer, Path initialDataPath) {

    public static final int DEFAULT_PLANT_COUNT = 200;
    public static final String DEFAULT_PRODUCER = "Chiquita";
    public static final Path DEFAULT_DATA_PATH = Path.of("/etc/os-release");

    public static @Template Configuration defaultConfiguration() {
        return new Configuration(DEFAULT_PLANT_COUNT, DEFAULT_PRODUCER, DEFAULT_DATA_PATH);
    }
}
```

See `examples/` for a demonstration of each possible style of configuration definitions:

- `examples/foo-missionmodel`: uses standard `@Parameter` configuration annotations.
- `examples/banananation`: (shown above) uses the `@Template` annotation to define a default `Configuration` object.
- `examples/config-with-defaults`: uses `@WithDefaults` to define a default for each parameter.
- `examples/config-without-defaults`: defined with no default arguments, requires all arguments to be supplied by the planner.

Use

The mission model may use a configuration to set initial values, for example:

```
this.sink = new Accumulator(0.0, config.sinkRate);
```

Activities

- [Activity Annotation](#)
- [Activity Metadata](#)
- [Activity Parameters](#)
- [Validations](#)
- [Activity Effect Model](#)
 - [A Note about Decomposition](#)
- [Computed Attributes](#)
- [Duration Types](#)

An activity is a modeled unit of mission system behavior to be utilized for simulations during planning. Activities emit events which can have various effects on modeled resources, and these effects can be modulated through input parameters. Activities can represent any part of the mission system behavior such as the availability of ground assets, or the execution of commands by the flight software or the instruments. In other words, activities in Merlin are entities whose role is to emit stimuli (events) to which the mission model reacts. Activities can therefore describe the relation: "when this activity occurs, this kind of thing should happen".

An activity type is a prototype for activity instances to be executed in a simulation. Activity types are defined by java classes that provide an `EffectModel` method to Merlin, along with a set of parameters. Each activity type exists in its own .java file, though activity types can be organized into hierarchical packages, for example as `gov.nasa.jpl.europa.clipper.gnc.TCMActivity`

Activity types consist of:

- metadata
- parameters that describe the range in execution and effects of the activity
- effect model that describes how the system will be perturbed when the activity is executed.

Activity Annotation

In order for Merlin to detect an activity type, its class must be annotated with the `@ActivityType` tag. An activity type is declared with its name using the following annotation:

```
@ActivityType("TurnInstrumentOff")
```

By doing so, the Merlin annotation processor can discover all activity types declared in the mission model, and validate that activity type names are unique.

Activity Metadata

Metadata of activities are structured such that the Merlin annotation processor can extract this metadata given particular keywords. Currently, the Merlin annotation processor recognizes the following tags: `contact`, `subsystem`, `brief_description`, and `verbose_description`.

These metadata tags are placed in a JavaDocs style comment block above the Activity Type to which they refer. For example:

```
/**
 * @subsystem Data
 * @contact mkumar
 * @brief_description A data management activity that deletes old files
 */
```

These tags are processed, at compile time, by the annotation processor to create documentation for the Activity types that are described in the mission model.

Activity Parameters

Activity parameters provide the ability to modulate the behavior of an activity instance's effect model. Aside from determining the effects of the

activity, these parameters can be used to determine its duration, decomposition into children activities and expansion into commands.

The Merlin annotation processor is used to extract and generate serialization code for parameters of activity types. The annotation processor also allows authors of a mission model to create mission-specific parameter types, ensuring that they will be recognized by the Merlin framework. See [Declaring Parameters](#) for a thorough explanation of all possible styles of parameter declarations. For more information on mission-specific parameter types, see [Value Mappers](#).

Validations

See [Declaring Parameters](#) for a thorough explanation of parameter validation.

Activity Effect Model

Every activity type has an associated "effect model" that describes how that activity impacts mission resources. An effect model is a method on the activity type class annotated with `@EffectModel`; there can only be one such method, and by convention it is named `run` that accepts the top-level `Mission` model as a parameter. This method is invoked when the activity begins, paused when the activity waits a period of time, and resumes when that period of time passes. The activity's computed duration will be measured from the instant this method is entered, and extends either until the method returns or until all spawned children of the activity have completed -- whichever is longer.

Exceptions: If an uncaught exception is thrown from an activity's effect model, the simulation will halt. No later activities will be performed, and simulation time will not proceed beyond the instant at which the exception occurred. As such, uncaught exceptions are essentially treated as fatal errors. We advise mission models to employ uncaught exceptions sparingly, as it deprives planners of information about the behavior of the spacecraft after the fault. (At least as of 2021-06-11, an uncaught exception will cause the simulation to abort without producing *any* results -- not even up to the point of failure.) On the other hand, uncaught exceptions may be useful to identify bugs in the mission model or situations that the mission model is not intended to simulate.

Actions: An activity's effect model may wait for time to pass, spawn other activities, and affect spacecraft state. Spacecraft state is affected via the `Mission` model parameter, which depends on the details of the modeled mission systems. (See [Developing a Mission Model](#) for more on mission modeling.)

Actions related to the passage of simulation time are provided as static methods on the `merlin.framework.ModelActions` class:

- `delay(duration)` : Delay the currently-running activity for the given duration. On resumption, it will observe effects caused by other activities over the intervening timespan.
- `waitFor(activityId)` : Delay the currently-running activity until the activity with specified ID has completed. On resumption, it will observe effects caused by other activities over the intervening timespan.
- `waitUntil(condition)` : Delay the currently-running activity until the provided `Condition` becomes true. On resumption, it will observe effects caused by other activities over the intervening timespan.

Actions related to spawning other activities are provided by the generated `ActivityActions` class, usually found under the `generated` package within your codebase.

- `spawn(activity)` : Spawn a new activity as a child of the currently-running activity at the current point in time. The child will initially see any effects caused by its parent up to this point. The parent will continue execution uninterrupted, and will not initially see any effects caused by its child.
- `call(activity)` : Spawn a new activity as a child of the currently-running activity at the current point in time. The child will initially see any effects caused by its parent up to this point. The parent will halt execution until the child activity has completed. This is equivalent to calling `waitFor(spawn(activity))`.

For example, consider a simple activity for running the on-board heaters:

```

@ActivityType("RunHeater")
public final class RunHeater {
    private static final int energyConsumptionRate = 1000;

    @Parameter
    public long durationInSeconds;

    @Validation("duration must be positive")
    public boolean validateDuration() {
        return durationInSeconds > 0;
    }

    @EffectModel
    public void run(final Mission mission) {
        spawn(new PowerOnHeater());

        final double totalEnergyUsed = durationInSeconds * energyConsumptionRate;
        mission.batteryCapacity.use(totalEnergyUsed);

        delay(durationInSeconds, Duration.SECONDS);

        call(new PowerOffHeater());
    }
}

```

This activity first spawns a `PowerOnHeater` activity, which then continues concurrently with the current `RunHeater` activity. Next, the total energy to be used by the heater is subtracted from the remaining battery capacity (see [Models and Resources](#)). The energy used depends on the duration parameter of the activity, allowing the activity's effect to be tuned by the planner. Next, the activity waits for the desired heater runtime to elapse, then spawns and waits for a `PowerOffHeater` activity. The activity completes after both children have completed.

A Note about Decomposition

In Merlin mission models, decomposition of an activity is not an independent method, rather it is defined within the effect model by means of invoking child activities. These activities can be invoked using the `call()` method, where the rest of the effect model waits for the child activity to complete; or using the `spawn()` method, where the effect model continues to execute without waiting for the child activity to complete. This method allows any arbitrary serial and parallel arrangement of child activities. This approach replaces duration estimate based wait calls with event based waits. Hence, this allows for not keeping track of estimated durations of activities, while also improving the readability of the activity procedure as a linear sequence of events.

Computed Attributes

Upon the termination of an Activity, the effect model can optionally log some information that will be included with the simulation results. This information has no effect on the current simulation, but can be used by downstream tools that process the simulation results.

To specify the information that will be logged, the mission modeler must change the return type of the method marked with the `@EffectModel` annotation.

Let's consider the `RunHeater` example from above:

```

@ActivityType("RunHeater")
public final class RunHeater {
    private static final int energyConsumptionRate = 1000;

    @Parameter
    public long durationInSeconds;

    @Validation("duration must be positive")
    public boolean validateDuration() {
        return durationInSeconds > 0;
    }

    @EffectModel
    public void run(final Mission mission) {
        spawn(new PowerOnHeater());

        final double totalEnergyUsed = durationInSeconds * energyConsumptionRate;
        mission.batteryCapacity.use(totalEnergyUsed);

        delay(durationInSeconds, Duration.SECONDS);

        call(new PowerOffHeater());
    }
}

```

The effect model return type is "void", which means no useful information will be logged. Let's change that to `String` and add a log message:

```
@EffectModel
public String run(final Mission mission) {
    spawn(new PowerOnHeater());

    final double totalEnergyUsed = durationInSeconds * energyConsumptionRate;
    mission.batteryCapacity.use(totalEnergyUsed);

    delay(durationInSeconds, Duration.SECONDS);

    call(new PowerOffHeater());

    return "This is a log message!"
}
}
```

Computed attributes are not limited to strings - they can be any type that merlin knows how to serialize. See [ValueMappers](#) to learn about what types are supported out of the box, and how to define custom serializers for your own types.

Let's show how we can define computed attributes as a Map of Strings to Longs:

```
@EffectModel
public Map<String, Long> run(final Mission mission) {
    spawn(new PowerOnHeater());

    final double totalEnergyUsed = durationInSeconds * energyConsumptionRate;
    mission.batteryCapacity.use(totalEnergyUsed);

    delay(durationInSeconds, Duration.SECONDS);

    call(new PowerOffHeater());

    return Map.of("duration_in_seconds", durationInSeconds,
        "energy_consumption_rate", (Long) energyConsumptionRate);
}
}
```

Duration Types

The scheduler (see [Scheduling Guide](#)) places activities in a plan to try achieving scheduling goals.

As the effect model of an activity is not accessible to the scheduler, its duration can only be computed by simulation. Satisfying temporal constraints associated with the scheduling of this activity ("activity A must end before activity B ") may lead to multiple simulations and thus more computation time.

However, it is possible to provide information about how the duration of an activity is determined to help the scheduler.

The duration of an activity is said to be:

- **controllable** if it is *only* determined by one of the activity parameter.
- **uncontrollable** otherwise.

By default, the duration of an activity is uncontrollable. To specify that the duration of an activity is controllable, the `@ControllableDuration` annotation can be added to the effect model method such as in the example below:

```

@ActivityType("RunHeater")
public final class RunHeater {
    private static final int energyConsumptionRate = 1000;

    @Parameter
    public long durationInSeconds;

    @Validation("duration must be positive")
    public boolean validateDuration() {
        return durationInSeconds > 0;
    }

    @EffectModel
    @ControllableDuration(parameterName = "durationInSeconds")
    public void run(final Mission mission) {
        spawn(new PowerOnHeater());

        final double totalEnergyUsed = durationInSeconds * energyConsumptionRate;
        mission.batteryCapacity.use(totalEnergyUsed);

        delay(durationInSeconds, Duration.SECONDS);
    }
}

```

Note that compared to the `RunHeater` activity present in the previous paragraphs, the `call(new PowerOffHeater())` at the end of the effect model has been removed because (1) it would have made the duration of the activity depend on the duration of the `PowerOffHeater` activity, (2) we assume `PowerOffHeater` has an uncontrollable duration. Spawning another activity does not affect the duration of this activity as they run in parallel.

However, let's say `PowerOffHeater` also has a controllable duration that is passed through a parameter, we could write the following activity:

```

@ActivityType("RunHeater")
public final class RunHeater {
    private static final int energyConsumptionRate = 1000;

    @Parameter
    public long durationInSeconds;

    @Validation("duration must be positive")
    public boolean validateDuration() {
        return durationInSeconds > 0;
    }

    @EffectModel
    @ControllableDuration(parameterName = "durationInSeconds")
    public void run(final Mission mission) {
        spawn(new PowerOnHeater());

        final double totalEnergyUsed = durationInSeconds * energyConsumptionRate;
        mission.batteryCapacity.use(totalEnergyUsed);

        durationPowerOff = 100L;
        durationPowerOn = durationInSeconds - durationPowerOff;

        delay(durationPowerOn, Duration.SECONDS);
        call(new PowerOffHeater(durationPowerOff))
    }
}

```

Even though the duration of `RunHeater` depends on the (fixed) duration of `PowerOffHeater`, we know its duration will be equal to `durationInSeconds`, making the activity effectively controllable.

The annotation `@ControllableDuration(parameterName = "durationInSeconds")` has no other effect than to tell the scheduler that the duration of this activity can be controlled via the `durationInSeconds` parameter. It acts like a contract between the mission model and the scheduler, ensuring that the duration of the activity will be equal to the duration specified by `durationInSeconds`.

After being given this information, the scheduler considers it can control the duration of the activity which will thus require less simulations, significantly improving scheduling performance for this activity type.

Activity Mappers

What is an Activity Mapper

An Activity Mapper is a Java class that implements the `ActivityMapper` interface for the `ActivityType` being mapped. It is required that each Activity Type in an mission model have an associated Activity Mapper, to provide several capabilities surrounding serialization/deserialization of activity instances.

The Merlin annotation processor can automatically [generate activity mappers](#) for every activity type, even for those with custom-typed parameters, but if it is desirable to create a custom activity mapper the interface is described below.

ActivityMapper Interface

The `ActivityMapper` interface is shown below:

```
public interface ActivityMapper<Instance> {
    String getName();
    Map<String, ValueSchema> getParameters();
    Map<String, SerializedValue> getArguments(Instance activity);

    Instance instantiateDefault();
    Instance instantiate(Map<String, SerializedValue> arguments) throws TaskSpecType.UnconstructableTaskSpecException;

    List<String> getValidationFailures(Instance activity);
}
```

The first thing to notice is that the interface takes a type parameter (here called `Instance`). When implementing the `ActivityMapper` interface, an activity mapper must supply the `ActivityType` being mapped. With that in mind, each of the methods shown must be implemented as such:

- `getName()` returns the name of the activity type being mapped
- `getParameters()` provides the named parameter fields of the activity along with their corresponding `ValueSchema`, that describes their structure
- `getArguments(Instance activity)` provides the actual values for each parameter from a provided activity instance
- `instantiateDefault()` creates a default instance of the activity type without any values provided externally
- `instantiate(Map<String, SerializedValue> arguments)` constructs an instance of the activity type from a the provided arguments, if possible
- `getValidationFailures(Instance activity)` provides a list of reasons a constructed activity is invalid, if any. Note that validation failures are different from instantiation errors. Validation failures occur when a constructed activity instance's parameters are outside acceptable range.

The `getParameters()` method returns a `Map<String, ValueSchema>`. In this map should be a key for every parameter, with a `ValueSchema` describing the structure of that parameter. See our [Value Schema documentation](#) for more information on creating value schemas.

Generated Activity Mappers

In most cases, you will likely want to let Merlin generate activity mappers for you. Thankfully, this is done automatically when running the Merlin Annotation Processor. When compiling your code with the Merlin annotation processor, the processor will produce an activity mapper for each activity type. This is made possible by the use of the `@WithMappers()` annotations in your `package-info.java`. Each java-file specified by these annotations is parsed to determine what types of values can be mapped. As long as there is a mapper for each activity parameter type used in the model, the annotation processor should have no issues creating activity mappers.

Value Mappers

Regardless of whether you create custom activity mappers or let Merlin generate them for you, you will likely find the need to work with a `ValueMapper` at some point. In fact, generating activity mappers is made quite simple by considering the fact that an activity instance is wholly defined by its parameter values.

You may find yourself asking "Just what _is_ a value mapper?" A value mapper is a small, focused class whose sole responsibility is to tell Merlin how to handle a specific type of value. Value mappers allow all sorts of capabilities from custom-typed activity parameters to custom-typed resources.

One of the most convenient things about using value mappers is the fact that Merlin comes with them already defined for all basic types. Furthermore, value mappers for combinations of types can easily be created by passing one `ValueMapper` into another during instantiation.

Although we provide value mappers for basic types, it is entirely acceptable to create custom value mappers for other types, such as those imported from external libraries. This can be done by writing a Java class which implements the `ValueMapper` interface. Below is a value mapper for an apache `Vector3D` type as an example:

```
public class Vector3DValueMapper implements ValueMapper<Vector3D> {

    @Override
    public ValueSchema getValueSchema() {
        return ValueSchema.ofSequence(ValueSchema.REAL);
    }

    @Override
    public Result<Vector3D, String> deserializeValue(final SerializedValue serializedValue) {
        return serializedValue
            .asList()
            .map(Result::<List<SerializedValue>, String>success)
            .orElseGet(() -> Result.failure("Expected list, got " + serializedValue.toString()))
            .match(
                serializedElements -> {
                    if (serializedElements.size() != 3) return Result.failure("Expected 3 components, got " + serializedElements.size());
                    final var components = new double[3];
                    final var mapper = new DoubleValueMapper();
                    for (int i=0; i<3; i++) {
                        final var result = mapper.deserializeValue(serializedElements.get(i));
                        if (result.getKind() == Result.Kind.Failure) return result.mapSuccess(_left -> null);

                        // SAFETY: `result` must be a Success variant.
                        components[i] = result.getSuccessOrNull();
                    }
                    return Result.success(new Vector3D(components));
                },
                Result::failure
            );
    }

    @Override
    public SerializedValue serializeValue(final Vector3D value) {
        return SerializedValue.of(
            List.of(
                SerializedValue.of(value.getX()),
                SerializedValue.of(value.getY()),
                SerializedValue.of(value.getZ())
            )
        );
    }
}
```

Notice there are just 3 methods to implement for a `ValueMapper`. The first is `getValueSchema()`, which should return a `ValueSchema` describing the structure of the value being mapped (see [here](#) for more info)

The next two methods are inverses of each other: `deserializeValue()` and `serializeValue()`. It is the job of `deserializeValue()` to take a `SerializedValue` and map it, if possible, into the mapper's supported value. Meanwhile, `serializeValue()` takes an instance of the mapper's supported value and turns it into a `SerializedValue`.

There are plenty of examples of value mappers over in the [contrib module](#).

Registering Value Mappers

As mentioned above, the `@WithMappers()` annotation is used to register value mappers for a mission model. Value mappers are expected to be defined with static constructor methods within classes listed in `@WithMappers()` annotations. For example, if `package-info.java` contains:

```
@WithMappers(BananaValueMappers.class)
```

Then the value mapper may define a custom `Configuration` value mapper with:

```
public final class BananaValueMappers {
    public static ValueMapper<Configuration> configuration() {
        return new ConfigurationValueMapper();
    }
}
```

Value mappers may be created for types that use parameterized types, but the parameterized types themselves must be either unbounded bounded or `Enum<>`. For example:

```
@Parameter
public List<? extends Foo> test;
```

or

```
@Parameter
public List<? extends Map<? super Foo, ? extends Bar>> test;
```

are not trivially resolved to a single value mapper due to the type constraints at play here.

What is a `SerializedValue`

When working with a `ValueMapper` it is inevitable that you will come across the `SerializedValue` type. This is the type we use for serializing all values that need serialization, such as activity parameters and resource values. In crafting a value mapper, you will have to both create a `SerializedValue` and parse one.

Constructing a `SerializedValue` tends to be more straightforward, because there are no questions about the structure of the value you are starting with. For basic types, you need only call `SerializedValue.of(value)` and the `SerializedValue` class will handle the rest. This can be done for values of the following types: `long`, `double`, `String`, `boolean`. Note that integers and floats can be represented by `long` and `double` respectively. For more complex types, you can also provide a `List<SerializedValue>` or `Map<String, SerializedValue>` to `SerializedValue.of()`. It is clear that these can be used to serialize lists and maps themselves, but arbitrarily complex structures can be serialized in this way. Consider the following examples:

```
int exInt = 5;
SerializedValue serializedInt = SerializedValue.of(exInt);

List<String> exList = List.of("a", "b", "c")
SerializedValue serializedList = SerializedValue.of(
    List.of(
        SerializedValue.of(exList.get(0)),
        SerializedValue.of(exList.get(1)),
        SerializedValue.of(exList.get(2))
    )
);

Map<String, Boolean> exMap = Map.of(
    "key1", true,
    "key2", false,
    "key3", true
);
SerializedValue serializedMap = SerializedValue.of(
    Map.of(
        "key1", SerializedValue.of(exMap.get("key1")),
        "key2", SerializedValue.of(exMap.get("key2")),
        "key3", SerializedValue.of(exMap.get("key3"))
    )
);

Vector3D exampleVec = new Vector3D(0,0,0);

SerializedValue serializedVec1 = SerializedValue.of(
    List.of(
        SerializedValue.of(exampleVec.getX()),
        SerializedValue.of(exampleVec.getY()),
        SerializedValue.of(exampleVec.getZ())
    )
);

SerializedValue serializedVec2 = SerializedValue.of(
    Map.of(
        "x", SerializedValue.of(exampleVec.getX()),
        "y", SerializedValue.of(exampleVec.getY()),
        "z", SerializedValue.of(exampleVec.getZ())
    )
);
```

The first 3 examples here are straightforward mappings from their java type to their serialized form, however the vector example is more interesting. To highlight this, two forms of `SerializedValue` have been given for it. In the first case, we serialize the `Vector3D` as a list of three values. This will work fine as long as whoever deserializes it knows that the list contains each component in order of x, y and z. In the second example, however, the vector is serialized as a map. Either of these representations may fit better in different scenarios. Generally, the structure of a `SerializedValue` constructed by a `ValueMapper` should match the `ValueSchema` the `ValueMapper` provides.

Example Activity Mapper

Below is an example of an Activity Type and its Activity mapper for reference:

Activity Type

```
@ActivityType("foo")
public final class FooActivity {
    @Parameter
    public int x = 0;

    @Parameter
    public String y = "test";

    @Parameter
    public List<Vector3D> vecs = List.of(new Vector3D(0.0, 0.0, 0.0));

    @Validation("x cannot be exactly 99")
    public boolean validateX() {
        return (x != 99);
    }

    @Validation("y cannot be 'bad'")
    public boolean validateY() {
        return !y.equals("bad");
    }

    @EffectModel
    public void run(final Mission mission) {
        // ...
    }
}
```

Generated Activity Mapper Example

```
package gov.nasa.jpl.aerie.foomissionmodel.generated.activities;

import gov.nasa.jpl.aerie.contrib.serialization.mappers.EnumValueMapper;
import gov.nasa.jpl.aerie.contrib.serialization.mappers.NullableValueMapper;
import gov.nasa.jpl.aerie.contrib.serialization.rulesets.BasicValueMappers;
import gov.nasa.jpl.aerie.foomissionmodel.Mission;
import gov.nasa.jpl.aerie.foomissionmodel.activities.FooActivity;
import gov.nasa.jpl.aerie.foomissionmodel.generated.ActivityTypes;
import gov.nasa.jpl.aerie.foomissionmodel.mappers.FooValueMappers;
import gov.nasa.jpl.aerie.merlin.framework.ActivityMapper;
import gov.nasa.jpl.aerie.merlin.framework.ModelActions;
import gov.nasa.jpl.aerie.merlin.framework.RootModel;
import gov.nasa.jpl.aerie.merlin.framework.ValueMapper;
import gov.nasa.jpl.aerie.merlin.protocol.driver.Topic;
import gov.nasa.jpl.aerie.merlin.protocol.model.Task;
import gov.nasa.jpl.aerie.merlin.protocol.types.InvalidArgumentsException;
import gov.nasa.jpl.aerie.merlin.protocol.types.Parameter;
import gov.nasa.jpl.aerie.merlin.protocol.types.SerializedValue;
import gov.nasa.jpl.aerie.merlin.protocol.types.UnconstructableArgumentException;
import gov.nasa.jpl.aerie.merlin.protocol.types.Unit;
import gov.nasa.jpl.aerie.merlin.protocol.types.ValueSchema;
import java.util.ArrayList;
import java.util.HashMap;
import java.util.List;
import java.util.Map;
import java.util.Optional;
import javax.annotation.processing.Generated;
import org.apache.commons.math3.geometry.euclidean.threed.Vector3D;

@Generated("gov.nasa.jpl.aerie.merlin.processor.MissionModelProcessor")
public final class FooActivityMapper implements ActivityMapper<RootModel<ActivityTypes, Mission>, FooActivity, Unit> {
    private final ValueMapper<Integer> mapper_x;

    private final ValueMapper<String> mapper_y;

    private final ValueMapper<List<Vector3D>> mapper_vecs;

    private final ValueMapper<Unit> computedAttributesValueMapper;

    private final Topic<FooActivity> inputTopic = new Topic<>();

    private final Topic<Unit> outputTopic = new Topic<>();

    @SuppressWarnings("unchecked")
    public FooActivityMapper() {
        this.mapper_x =
            BasicValueMappers.$int();
        this.mapper_y =
            new NullableValueMapper<>(
                BasicValueMappers.string());
    }
}
```

```

    this.mapper_vecs =
        new NullableValueMapper<>{
            BasicValueMappers.list(
                FooValueMappers.vector3d(
                    BasicValueMappers.$double()));
    this.computedAttributesValueMapper =
        new EnumValueMapper(Unit.class);
}

@Override
public List<String> getRequiredParameters() {
    return List.of();
}

@Override
public ArrayList<Parameter> getParameters() {
    final var parameters = new ArrayList<Parameter>();
    parameters.add(new Parameter("x", this.mapper_x.getValueSchema()));
    parameters.add(new Parameter("y", this.mapper_y.getValueSchema()));
    parameters.add(new Parameter("vecs", this.mapper_vecs.getValueSchema()));
    return parameters;
}

@Override
public Map<String, SerializedValue> getArguments(final FooActivity activity) {
    final var arguments = new HashMap<String, SerializedValue>();
    arguments.put("x", this.mapper_x.serializeValue(activity.x));
    arguments.put("y", this.mapper_y.serializeValue(activity.y));
    arguments.put("vecs", this.mapper_vecs.serializeValue(activity.vecs));
    return arguments;
}

@Override
public FooActivity instantiate(final Map<String, SerializedValue> arguments) throws
    InvalidArgumentsException {
    final var template = new FooActivity();
    Optional<Integer> x = Optional.ofNullable(template.x);
    Optional<String> y = Optional.ofNullable(template.y);
    Optional<List<Vector3D>> vecs = Optional.ofNullable(template.vecs);

    final var invalidArgsExBuilder = new InvalidArgumentsException.Builder("activity", "foo");

    for (final var entry : arguments.entrySet()) {
        try {
            switch (entry.getKey()) {
                case "x":
                    template.x = this.mapper_x.deserializeValue(entry.getValue())
                        .getSuccessOrThrow(failure -> new UnconstructableArgumentException("x", failure));
                    break;
                case "y":
                    template.y = this.mapper_y.deserializeValue(entry.getValue())
                        .getSuccessOrThrow(failure -> new UnconstructableArgumentException("y", failure));
                    break;
                case "vecs":
                    template.vecs = this.mapper_vecs.deserializeValue(entry.getValue())
                        .getSuccessOrThrow(failure -> new UnconstructableArgumentException("vecs", failure));
                    break;
                default:
                    invalidArgsExBuilder.withExtraneousArgument(entry.getKey());
            }
        } catch (final UnconstructableArgumentException e) {
            invalidArgsExBuilder.withUnconstructableArgument(e.parameterName, e.failure);
        }
    }

    x.ifPresentOrElse(
        value -> invalidArgsExBuilder.withValidArgument("x", this.mapper_x.serializeValue(value)),
        () -> invalidArgsExBuilder.withMissingArgument("x", this.mapper_x.getValueSchema()));
    y.ifPresentOrElse(
        value -> invalidArgsExBuilder.withValidArgument("y", this.mapper_y.serializeValue(value)),
        () -> invalidArgsExBuilder.withMissingArgument("y", this.mapper_y.getValueSchema()));
    vecs.ifPresentOrElse(
        value -> invalidArgsExBuilder.withValidArgument("vecs", this.mapper_vecs.serializeValue(value)),
        () -> invalidArgsExBuilder.withMissingArgument("vecs", this.mapper_vecs.getValueSchema()));

    invalidArgsExBuilder.throwIfAny();
    return template;
}

@Override
public List<String> getValidationFailures(final FooActivity activity) {
    final var failures = new ArrayList<String>();
    if (!activity.validateX()) failures.add("x cannot be exactly 99");
    if (!activity.validateY()) failures.add("y cannot be 'bad'");
    return failures;
}

@Override

```

```

~
public ValueSchema getReturnValueSchema() {
    return this.computedAttributesValueMapper.getValueSchema();
}

@Override
public SerializedValue serializeReturnValue(final Unit returnValue) {
    return this.computedAttributesValueMapper.serializeValue(returnValue);
}

public Topic<FooActivity> getInputTopic() {
    return this.inputTopic;
}

public Topic<Unit> getOutputTopic() {
    return this.outputTopic;
}

@Override
public Task<Unit> createTask(final RootModel<ActivityTypes, Mission> model,
    final FooActivity activity) {
    return ModelActions
        .threaded(() -> {
            try (final var restore = model.registry().contextualizeModel(model)) {
                ModelActions.emit(activity, this.inputTopic);
                activity.run(model.model());
                ModelActions.emit(Unit.UNIT, this.outputTopic);
                return Unit.UNIT;
            }
        })
        .create(model.executor());
}
}

```

Value Schemas

A Value What?

A value *schema*! A value schema is a description of the structure of some value. Using value schemas, users can tell our system how to work with arbitrarily complex types of values, so long as they can be described using the value schema constructs provided by Merlin. Let's take a look at how it's done.

Starting With The Basics

At a fundamental level, a value schema is no more than a combination of elementary value schemas. Merlin defines the elementary value schemas, so let's take a look at them:

- `REAL` : A real number
- `INT` : An integer
- `BOOLEAN` : A boolean value
- `STRING` : A string of characters
- `DURATION` : A duration value
- `PATH` : A file path
- `VARIANT` : A string value constrained to a set of acceptable values.

If you are trying to write a value schema for an integer value, all you have to do is use the `INT` value schema, but of course values can quickly take on more complex structures, and for that we must examine the remaining value schema constructs.

A Note About Variants

The `Variant` value schema is a little unique among the elementary value schemas in that it requires input, the set of acceptable values. The way to provide this set of values depends on the context in which you are creating a value schema and will be addressed in the corresponding section below.

Building Things Up

In order to combine elementary value schemas, we provide two main constructs:

- `SERIES` : Denotes a list of values of a single type
- `STRUCT` : Denotes a structure of independent values of varying types

The `SERIES` node allows a straightforward declaration of a list of values that fall under the same schema, while the `STRUCT` node opens things up, allowing you to create any combination of different values, each labeled by some string name.

Now that you've seen the basics, let's talk about the two different ways to create value schemas -- in code and in JSON/GraphQL (serialized value schemas).

Value Schemas From Code

Creating a value schema from code is straightforward thanks to the `ValueSchema` class. Each of the elementary value schemas is accessible as via the `ValueSchema` class. For example, a `REAL` is given by `ValueSchema.REAL`. The one exception is that to create a `VARIANT` type value schema you'll need to call `ValueSchema.ofVariant(Class<? extends Enum> enum)`, providing an enum with to specify the acceptable variants.

Like the `VARIANT` element, the `SERIES` and `STRUCT` constructs are created by calling their corresponding methods `ValueSchema.ofSeries(ValueSchema value)` and `ValueSchema.ofStruct(Map<String, ValueSchema> map)`.

Examples Using `ValueSchema`

Below are a few examples of how to create a `ValueSchema`. In each, a Java type and its corresponding `ValueSchema` are compared.

- `Integer` is described by `ValueSchema.INT`

- `List<Double>` is described by `ValueSchema.ofSeries(ValueSchema.REAL)`
- `Float[]` is described by `ValueSchema.ofSeries(ValueSchema.REAL)`

Note that the second and third examples are entirely different Java types, but are represented by the same `ValueSchema`. It is also important to take a look at a `Map` type, as it can be confusing at first how to represent its structure:

`Map<String, Integer>` is described by

```
ValueSchema.ofStruct(
    Map.of(
        "keys": ValueSchema.ofSeries(ValueSchema.STRING),
        "values": ValueSchema.ofSeries(ValueSchema.INT)
    )
)
```

Here we are taking note of the fact that a `Map` is really just a list of keys and a list of values. As a final example, consider the custom type below:

```
public class CustomType {
    public int foo;
    public boolean bar;
    public List<String> baz
}
```

A variable of type `CustomType` has structure described by:

```
ValueSchema.ofStruct(
    Map.of(
        "foo": ValueSchema.INT,
        "bar": ValueSchema.BOOLEAN,
        "baz": ValueSchema.ofSeries(ValueSchema.STRING)
    )
)
```

Serialized Value Schemas

Creating value schemas from JSON/GraphQL is a little less straightforward, since your IDE won't be able to help you, but fear not, you've come to the right place. A value schema is created by declaring an object with a `type` field that tells which type of schema is being created. The values allowed in this field are given below:

- `"real"` corresponds to `REAL`
- `"int"` corresponds to `INT`
- `"boolean"` corresponds to `BOOLEAN`
- `"string"` corresponds to `STRING`
- `"duration"` corresponds to `DURATION`
- `"path"` corresponds to `PATH`
- `"variant"` corresponds to `VARIANT`
- `"series"` corresponds to `SERIES`
- `"struct"` corresponds to `STRUCT`

Variant

For the `"variant"` type, you'll need to include a second field called `variants` whose value is a list of objects specifying the string-valued `key` and `label` fields of each variant like this:

```
{
  "type": "variant",
  "variants": [
    {
      "key": "ON",
      "label": "ON"
    },
    {
      "key": "OFF",
      "label": "OFF"
    }
  ]
}
```

Series

For the `"series"` type, a second field called `"items"` must be included as well that provides the value schema for the items in the series. See the below example, a value schema for a list of integers.

```
{
  "type": "series",
  "items": {
    "type": "int"
  }
}
```

Struct

Lastly, for the `"struct"` type, a second field called `"items"` must be included that provides the actual structure of the struct, mapping string keys to their corresponding value schema. See the below example, a value schema for a struct with a string-valued `label` field, real-valued `position` field, and boolean-valued `on` field:

```
{
  "type": "struct",
  "items": {
    "label": { "type": "string" },
    "position": { "type": "real" },
    "on": { "type": "boolean" }
  }
}
```

Examples Creating Serialized Value Schemas

Below are more examples of creating serialized value schemas using JSON:

1. A value schema for an integer:

```
{
  "type": "int"
}
```

2. A value schema for a list of paths:

```
{
  "type": "series",
  "items": {
    "type": "path"
  }
}
```

3. A value schema for a list of lists of booleans:

```
{
  "type": "series",
  "items": {
    "type": "series",
    "items": {
      "type": "boolean"
    }
  }
}
```


4. A value schema for a structure containing a list of integers labeled `lints`, and a boolean labeled `active`:

```
{
  "type": "struct",
  "items": {
    "lints": {
      "type": "series",
      "items": { "type": "int" }
    },
    "active": {
      "type": "boolean"
    }
  }
}
```

Declaring Parameters

The Merlin interface offers a variety of ways to define **parameters** for mission model configurations and activities. Parameters offer a concise way to export information across the mission-agnostic Merlin interface – namely a parameter's type to support serialization and a parameter's "required" status to ensure that parameters without mission-model-defined defaults have an argument supplied by the planner.

In this guide **parent class** refers to the Java class that encapsulates parameters. This class may take the form of either a mission model configuration or activity.

Both configurations and activities make use of the same Java annotations for declaring parameters within a parent class. The `@Export` annotation interface serves as the common qualifier for exporting information across the mission-agnostic Merlin interface. The following parameter annotations serve to assist with parameter declaration and validation:

- `@Export.Parameter``
- `@Export.Template``
- `@Export.WithDefaults``
- `@Export.Validation``

The following sections delve into each of these annotations along with examples and suggested use cases for each.

Declaration

Without Export Annotations

The first – and perhaps less obvious option – is to not use any parameter annotations. If a parent class contains no `@Export.Parameter``, `@Export.Template``, or `@Export.WithDefaults`` annotation it is assumed that every class member variable is a parameter to export to Merlin.

Defining a parent class becomes as simple as `public record Configuration(Integer a, Integer b, Integer c) { }`. However, it is not possible to declare a member variable that is not an exported parameter with this approach.

Example

```
@ActivityType("Simple")
public record SimpleActivity(Integer a, Integer b, Integer c) {

    @EffectModel
    public void run(final Mission mission) {
        mission.count.set(a);
        delay(1, SECOND);
        mission.count.set(b);
        delay(1, SECOND);
        mission.count.set(c);
    }
}
```

In the above example `a`, `b`, and `c` are all parameters that require planner-provided arguments at planning time.

See Also

Aerie's `examples/config-without-defaults` makes use of this succinct style for declaring mission model configuration parameters.

For more information on records in Java 16+, see [Java Record Classes](#).

Recommendation

Avoid `@Export.Parameter``, `@Export.Template``, or `@Export.WithDefaults`` when every member variable for a parent class should be an exported parameter without a default value.

This approach is great for defining a simple `record` type parent class without defaults. If defaults are desired then `@Template` or `@WithDefaults` can be used without changing a `record` type class definition.

`@Export.Parameter``

The `@Parameter` annotation is the most explicit way to define a parameter and its defaults. Explicitly declaring each parameter within a parent class with or without a default value gives the mission modeler the freedom to decide which member variables are parameters and which

parameters are required by the planner.

Example

```
public final class Configuration {

    public Integer a;

    @Export.Parameter
    public Integer b;

    @Export.Parameter
    public Integer c = 3;

    public Configuration(final Integer a, final Integer b, final Integer c) {
        this.a = a;
        this.b = b;
        this.c = c;
    }
}
```

In the above example the parent class is a mission model configuration. Here's a close look at each member variable:

- `a`: a traditional member variable. Not explicitly declared as a parameter and therefore is not considered to be a parameter.
- `b`: explicitly declared as a parameter without a default value. A value will be required by the planner.
- `c`: explicitly declared as a parameter with a default value. A value will not be required by the planner.

See Also

Aerie's `examples/foo-missionmodel` uses this style when declaring mission model configuration parameters.

Recommendation

Declare each parameter with a `@Parameter` when a non-`record` type parent class is desired.

Some mission modelers may prefer the explicitness provided by individual `@Parameter` annotations. However, this opens the door to subtle mistakes such as an unintentionally absent `@Parameter` annotation or an unintentionally absent default assignment. Those who prefer a more data-oriented approach may also find this style to not be as ergonomic as using a simple `record` type.

`@Export.Template`

The `@Template` annotation decouples parameter definitions and default values, allowing `record` types to be used as parent classes. When the `@Template` annotation is used every parent class member variable is interpreted as a parameter to export to Merlin. This annotation must be attached to a `public static` constructor method.

Example

```
@ActivityType("ThrowBanana")
public record ThrowBananaActivity(double speed) {

    @Template
    public static ThrowBananaActivity defaults() {
        return new ThrowBananaActivity(1.0);
    }

    @Validation("Speed must be positive")
    public boolean validateBiteSize() {
        return this.speed() > 0;
    }

    @EffectModel
    public void run(final Mission mission) {
        mission.plant.add(-1);
    }
}
```

In the above example `ThrowBananaActivity` is a `record` type with one constructor parameter, `speed`.

See Also

Aerie's `examples/banananation` uses this style within `GrowBananaActivity` and `ThrowBananaActivity`.

Recommendation

Use `@Template` when every member variable for a parent class should be an exported parameter with a default value.

@WithDefaults

Similarly to `@Template`, `@WithDefaults` annotation also decouples parameter definitions and default values, allowing `record` types to be used as parent classes. When the `@WithDefaults` annotation is used every parent class member variable is interpreted as a parameter to export to Merlin. Unlike `@Template`, a sparse set of default values may be supplied.

This annotation must be attached to a nested `public static class` within the parent class. Each member variable of this nested class must have the same name as a parent class's member variable. Not every parent class member variable is required to have an associated member variable within the nested class. This allows the mission modeler to selectively choose which parameters must be supplied by the planner.

Example

```
public record Configuration(Integer a, Double b, String c) {  
  
    @WithDefaults  
    public static final class Defaults {  
        public static Integer a = 42;  
        public static String c = "JPL";  
    }  
}
```

In the above example the parent class is a mission model configuration. Here's a close look at each member variable:

- `a`: a parameter with an associated default value.
- `b`: a parameter without a default value. A value will be required by the planner.
- `c`: a parameter with an associated default value.

See Also

Aerie's `examples/config-with-defaults` uses this style within its mission model configuration. The `examples/banananation` mission model also uses this style within `BakeBananaBreadActivity`.

Recommendation

Use `@WithDefaults` when every member variable for a parent class should be an exported parameter with an optionally provided default value.

Validation

A mission model configuration or activity instance can be validated by providing one or more methods annotated by `@Validation`. The annotation message specifies the message to present to a planner when the validation fails. For example:

```
@Validation("instrument power must be between 0.0 and 1000.0")  
public boolean validateInstrumentPower() {  
    return (instrumentPower_W >= 0.0) && (instrumentPower_W <= 1000.0);  
}
```

The Merlin annotation processor identifies these methods and arranges for them to be invoked whenever the planner instantiates an instance of the class. A message will be provided to the planner for each failing validation, so the order of validation methods does not matter.

Models & Resources

Mission Resources

In Merlin, a resource is any measurable quantity whose behavior is to be tracked over the course of a simulation. Resources are general-purpose, and can model quantities such as finite resources, geometric attributes, ground and flight events, and more. Merlin provides basic models for three types of quantity: a discrete quantity that can be set (`Register`), a continuous quantity that can be added to (`Counter`), and a continuous quantity that grows autonomously over time (`Accumulator`).

A common example of a `Register` would be a spacecraft or instrument mode, while common `Accumulator` s might be battery capacity or data volume.

Defining a resource is as simple as constructing a model of the appropriate type. The model will automatically register its resources for use from the Aerie UI. Alternatively, a resource may be **derived** or **sampled** from an existing resource.

Derived Resources

A derived resource is constructed from an existing resource given a mapping transformation.

For example, the `Imager` sample model defines an "imaging in progress" resource with:

```
this.imagingInProgress = this.imagerMode.map($ -> $ != ImagerMode.OFF);
```

In this example, `imagingInProgress` is a full-fledged discrete resource and will depend only on the imager's on/off state.

A derived resource may also be constructed from a real resource. For example, given `Accumulator` s `instrumentA` and `instrumentB` , a resource that maintains the current sum of both volumes may be constructed with:

```
var sumResource = instrumentA.volume.resource.plus(instrumentB.volume.resource);
```

Sampled Resources

A sampled resource allows for a new resource to be constructed from arbitrarily many existing resources/values and to be sampled once per second. This differs from a derived resource which provides a continuous mapping transformation from a single existing resource.

For example, the `Mission` sample model defines a "battery state of charge" resource with:

```
this.batterySoC = new SampledResource<>(() -> this.source.volume.get() - this.sink.volume.get());
```

In this example, `batterySoC` will be updated once per second to with the current difference between the "source" volume and "sink" volume.

Custom models

Often, the semantics of the pre-existing models are not exactly what you need in your adaptation. Perhaps you'd like to prevent activities from changing the rate of an `Accumulator` , or you'd like to have some helper methods for interrogating one or more resources. In these cases, a custom model may be a good solution.

A custom model is a regular Java class, extending the `Model` class generated for your adaptation by Merlin (or the base class provided by the framework, if it's mission-agnostic). It may implement any helper methods you'd like, and may contain any sub-models that contribute to its purpose. The only restriction is that it **must not** contain any mutable state of its own -- all mutable state must be held by one of the basic models, or one of the internal state-management entities they use, known as "cells".

The `contrib` package is a rich source of example models. See [the repository](#) for more details.

Creating Plans

Plans can be created via [Planning Web GUI](#) or interfacing with the [Aerie GraphQL API](#).

Constraints

- [Defining Constraints](#)
- [Creating A Constraint](#)
- [Constraint Examples](#)
- [Constraint Violation Examples](#)
- [Constraints API](#)

Overview

When analyzing a simulation's results, it may be useful to detect windows where certain conditions are met. Constraints are the Aerie tool for fulfilling that role. A constraint is a condition on activities and resources that must hold through an entire simulation. If a constraint does not hold true at any point in a simulation, this is considered a violation. The results yielded by a simulation run will include a list of violation windows for every constraint that has been violated.

Defining Constraints

Constraints are defined in with a typescript API via the Aerie GraphQL API or in the Aerie UI. To define a constraint, a Constraint object must be constructed and returned from an `export default function` (see our [comprehensive documentation of the constraints API](#)).

Creating a Constraint in Aerie

Constraints can be created via the Aerie GraphQL API. For an example of how to create a constraint for a mission model see the example, [Creating a constraint for a mission model](#) detailed in the [Aerie GraphQL API Software Interface Specification](#).

Constraint Examples

To define a constraint, you will need to build it up by constructing and transforming objects in the [constraint API](#). To help get you started, here are a few examples:

Constraint Example 1

Let's start off with a basic constraint that a resource, let's call it `BatteryTemperature`, doesn't exceed some threshold, say 340. We do so by using `Real.Resource(...)` to get the `BatteryTemperature` resource, and `Real.Value(...)` to get a real number we can compare a real resource profile to.

```
export default (): Constraint => Real.Resource("BatteryTemperature")
  .lessThanOrEqualTo(Real.Value(340))
```

The `Real.Resource(...)` function creates an object that refers to the `BatteryTemperature` real resource profile. The `.lessThanOrEqualTo(...)` method then expects either another real profile or a number literal as argument. In the above example, we passed it a real profile which has the value 340 for all time. We could instead omit `Real.Value(...)`, and `.lessThanOrEqualTo(...)` will automatically wrap the 340 literal in `Real.Value(...)` for us:

```
// This is identical
export default (): Constraint => Real.Resource("BatteryTemperature")
  .lessThanOrEqualTo(340)
```

The result of `.lessThanOrEqualTo(...)` is a `Windows` object, representing the time windows when the condition is true.

Constraint Example 2

Now we examine a more complex constraint. Let's imagine a solar panel that rotates the panels to a certain angle. Suppose the panels can rotate as fast as 5 degrees per second, but are not allowed to go more than 3 degrees per second unless the spacecraft is operating in IDLE mode. For this we will use a real resource, `PanelAngle`, and a discrete resource, `OpMode`.

Note that this breaks down to two conditions, either of which must be true the entire simulation. This constraint should be satisfied as either:

1. The `OpMode` is "IDLE"

2. The rate of the `PanelAngle` is no more than 3 degrees per second

```
export default (): Constraint => Windows.Any(
  Discrete.Resource("OpMode").equal("IDLE"),
  Real.Resource("PanelAngle").rate().lessThan(3)
)
```

The API keeps track of the type schemas of all your Discrete and Real value profiles. Real profiles are easy; they are always numbers. The structures of Discrete profiles are defined by the mission modeler. For example, the `OpMode` resource might be defined as an enum of either `"IDLE"` or `"ACTIVE"`. If you tried to use a different value, like `Discrete.Resource("OpMode").equal("BOOGIE")`, it would throw a compile-time type error.

Much like the previous example, the `.equal("IDLE")` method could instead be `.equal(Discrete.Value("IDLE"))`. The `equal` method and all such comparison operators operate on Discrete and Real profiles of the same type. If you provide a literal instead of a profile, it will be automatically wrapped in `Discrete.Value(...)` if it is the correct type.

Constraint Example 3

The first example of an activity constraint we present says that whenever an instance of `ActivityTypeA` occurs, the value of `ResourceX` must be less than 10.0.

```
export default(): Constraint => Constraint.ForEachActivity(
  ActivityType.ActivityTypeA,
  instance => Real.Resource("ResourceX").lessThan(10)
    .if(instance.window())
)
```

For those unfamiliar with Typescript, the `instance => ...` syntax defines an anonymous function which the `Constraint.ForEachActivity(...)` function calls. `instance` is of the type `ActivityInstance<A extends ActivityType>`, and can be used to access the instance's window, start time, end time, and parameters.

The `.if(instance.window())` is very important. Without it, the constraint would apply at all times, not just during `ActivityTypeA`. In general, `expressionWindows.if(conditionWindows)` means that `expressionWindows` is only checked during the windows of `conditionWindows`. It is equivalent to `Windows.Any(conditionWindows.invert(), expressionWindows)`, meaning that the above example could also be written:

```
export default(): Constraint => Constraint.ForEachActivity(
  ActivityType.ActivityTypeA,
  instance => Windows.Any(
    instance.window().invert(),
    Real.Resource("ResourceX").lessThan(10)
  )
)
```

Constraint Example 4

As a final example, we present a complex constraint containing two activity types, and several nested expressions. Most constraints should be much simpler than this, but this demonstrates just how capable the constraints API is.

The constraint below basically says this: For each pair of activities of type TypeA and TypeB, during the intersection of the two activities either parameter `b` of the TypeB instance must be false, or Resource ResC must be no greater than half of parameter `a` of the TypeA instance. Quite a mouthful, but here it is:

```
export default(): Constraint => Constraint.ForEachActivity(
  ActivityType.TypeA,
  instanceA => Constraint.ForEachActivity(
    ActivityType.TypeB,
    instanceB => Windows.Any(
      instanceB.parameters.b.equal(false),
      Real.Resource("ResC").lessThanOrEqual(instanceA.parameters.a.times(0.5))
    ).if(instanceA.window()).if(instanceB.window())
  )
)
```

Violation Examples

Constraint violations contain two sets of information describing where constraints are violated. First, a list of associated activity instance IDs representing the activity instances in violation (this will be an empty list for constraints that don't involve activities). Second, the list of violation windows themselves tells when during the simulation violations occur.

Constraint violations are reported per activity instance, so it is entirely possible for multiple violations to be produced by a single constraint. This unambiguous representation clearly indicates activity instances that violate a constraint despite the constraint being defined at the type-level.

Below are several examples of constraint violations:

A single activity instance with ID "2" is in violation from time 5 to 7:

```
{
  "activityInstanceIds": [ "2" ],
  "windows": [ [5, 7] ]
}
```

A constraint is violated from 2 to 4 and also from 5 to 8. No activities are involved in this violation.

```
{
  "activityInstanceIds": [],
  "windows": [ [2, 4], [5, 8] ]
}
```

Constraints API

Below we list all classes, constructors, and methods in the Constraints API, along with their arguments and return types. But before we do that, keep in mind that the entire API is a facade and does *none* of the constraint evaluation itself. So for example, when we say that `.lessThan(...)` returns a "set of windows", this is merely a helpful lie.

In reality, it returns a node of an Abstract Syntax Tree (AST) which represents the `LessThan` operation, which is serialized and sent to Merlin's Java backend to be evaluated there. Before the constraints API was released, constraint authors had to write a highly verbose JSON AST representing these operations. The API builds that same AST. Due to this architecture, it is not possible to use the constraint to actually *observe* the "set of windows" represented by `.lessThan(...)`, or the "discrete profile" represented by `Discrete.Resource(...)`, or any other value or object "returned" by the API.

Constraint

The `Constraint` class is what the constraint function returns, and can represent all possible constraints. In some use cases, the constraint author will directly create a `Constraint`, such as with the constructor functions below. In many other use cases, the author will return a `Windows` object instead, which will be automatically converted to a `Constraint` as if they had called the `windows.violations()` method.

constructor `ForEachActivity`

Evaluates a constraint on each instance of an activity type, and evaluates to the aggregated list of violations

- **Arguments:**
 - `activityType`: variant of `ActivityType`
 - `expression`: A function of signature `(instance: ActivityInstance<A extends ActivityType>) => Constraint`
- **Returns:** `Constraint`

constructor `ForbiddenActivityOverlap`

Evaluates to a list of violations, one for each overlap between the given activity types.

- **Arguments:**
 - `activityType1`: variant of `ActivityType`
 - `activityType2`: variant of `ActivityType`
- **Returns:** `Constraint`

ActivityType

An enum of activity type names. The variants are generated automatically from the mission model.

ActivityInstance<A extends ActivityType>

Represents an instance of an activity type. It is only intended to be created by the `ForEachActivity` constructor and accessed by the function

method `window`

Gets the window from start to end of the activity instance

- **Returns:** `Windows`

method `start`

Gets the instantaneous window at the start of the activity instance

- **Returns:** `Windows`

method `end`

Gets the instantaneous window at the end of the activity instance

- **Returns:** `Windows`

getter method `parameters`

Gets an object containing each of the activity parameters' referenced profile as a field

- **Returns:** `ActivityParameters`

Windows

Represents a set of time intervals (a.k.a. windows).

constructor `All`

Computes the intersection of the input windows

- **Arguments:**
 - `...windows` : any number of `Windows` objects
- **Returns:** `Windows`

constructor `Any`

Computes the union of the input windows

- **Arguments:**
 - `...windows` : any number of `Windows` objects
- **Returns:** `Windows`

method `if`

Transforms a windows object, filtering `this`'s windows by only checking them during `condition`'s windows.

`expressionWindows.if(conditionWindows)` is equivalent to `Windows.Any(conditionWindows.invert(), expressionWindows)`.

- **Arguments:**
 - `condition` : `Windows`
- **Returns:** `Windows`

method `invert`

Transforms a `windows` object, inverting the windows. This is the complement of the object with the plan bounds.

- **Returns:** `Windows`

method `shorterThan`

Transforms a `Windows` object, returning the windows whose duration is shorter than `duration`.

- **Arguments:**

- `duration` : `long`

- **Returns:** `Windows`

method `longerThan`

Transforms a `Windows` object, returning the windows whose duration is longer than `duration`.

- **Arguments:**

- `duration` : `long`

- **Returns:** `Windows`

method `shiftBy`

Transforms a `Windows` object, shifting the start and end of each window by a fixed positive or negative duration, respectively `fromStart` and `fromEnd`. Window [1,2] shifted by [-1,1] will become [0,3]. The same [1,2] shifted by [1,-2] will result in the `empty` window and disappear from the set of windows.

- **Arguments:**

- `fromStart` : `long`
- `fromEnd` : `long`

- **Returns:** `Windows`

method `violations`

Converts a `Windows` object into a list of violations, one for each interval that does *not* have a window. This method is always optional, and only exists if you want to be explicit about converting a `Windows` object to a `Constraint`. If you don't call this method, it will be done for you.

- **Returns:** `Constraint`

`Real`

A real number profile. This essentially represents a function from time to real numbers. You cannot sample or query the function from within the constraints API, but you can transform it and convert it into `Windows` when certain conditions are met.

constructor `Resource`

References the profile defined by a Real Resource. The string resource name is constrained to be a valid real resource name; hence the argument is `RealResourceName` instead of `string`. `RealResourceName` merely defines a small subset of all strings to be valid arguments.

- **Arguments:**

- `name` : `RealResourceName` name of the resource to get the profile of

- **Returns:** `Real`

constructor `Value`

Creates a constant profile which is equal to the argument for all time. This function is optional when used on the right side of an operator. For example, `Real.Resource("res").equal(Real.Value(3))` can be written as `Real.Resource("res").equal(3)`; but `Real.Value(3).equal(Real.Resource("res"))` *cannot* be written `3.equal(Real.Resource("res"))`.

- **Arguments:**

- `value` : Real number to build a profile from

- **Returns:** `Real`

method **plus**

Adds this real profile to another

- **Arguments:**

- **other**: Either a **Real** or a **number**

- **Returns:** **Real**

method **plus**

Multiplies this real profile by a constant

- **Arguments:**

- **multiplier**: **number**

- **Returns:** **Real**

method **plus**

Computes the derivative of the profile

- **Returns:** **Real**

method **lessThan**

Computes the windows when this is less than another real profile.

- **Arguments:**

- **other**: Either a **Real** or a **number**

- **Returns:** Set of windows

method **lessThanOrEqual**

Computes the windows when this is less than or equal to another real profile.

- **Arguments:**

- **other**: Either a **Real** or a **number**

- **Returns:** Set of windows

method **greaterThan**

Computes the windows when this is greater than another real profile.

- **Arguments:**

- **other**: Either a **Real** or a **number**

- **Returns:** Set of windows

method **greaterThanOrEqual**

Computes the windows when this is greater than or equal to another real profile.

- **Arguments:**

- **other**: Either a **Real** or a **number**

- **Returns:** Set of windows

method **changes**

Computes the instantaneous windows when the profile changes

- **Returns:** **Windows**

method `equal`

Computes the windows when this profile is equal to another

- **Arguments:**

- `other` : `Real`

- **Returns:** `Windows`

method `notEqual`

Computes the windows when this profile is not equal to another

- **Arguments:**

- `other` : `Real`

- **Returns:** `Windows`

`Discrete<Schema>`

A discrete value profile, where `Schema` is the type of the value. This essentially represents a function from time to *any* codomain (set of possible outputs), and the codomain is the `Schema` type. `Schema` is usually decided by the resource or activity parameter that you want to access, and the `Discrete` class imposes a type requirement that any comparisons must be between Discrete objects of the same `Schema` type.

You cannot sample or query the function from within the constraints API, but you can transform it and convert it into `Windows` when certain conditions are met.

constructor `Resource`

References the profile defined by a Discrete Resource. The string resource name is constrained to be a valid real resource name; hence the argument type is `R extends ResourceName` instead of `string`. `ResourceName` merely defines a small subset of all strings to be valid arguments. `DiscreteResourceSchema<R>` matches resource name to the discrete `Schema` (profile codomain).

- **Arguments:**

- `name` : `R extends ResourceName` ; name of the resource to get the profile of

- **Returns:** `Discrete<DiscreteResourceSchema<R>>`

constructor `Value`

Creates a constant profile which is equal to the argument for all time. This function is optional when used on the right side of an operator. For example, `Discrete.Resource("res").equal(Discrete.Value("value"))` can be written as `Discrete.Resource("res").equal("value")` ; but `Discrete.Value("value").equal(Discrete.Resource("res"))` *cannot* be written `"value".equal(Discrete.Resource("res"))` .

- **Arguments:**

- `value` : any type, which determines `Schema` ; Any value that can be serialized

- **Returns:** `Discrete<Schema>`

method `transition`

Returns the set of instantaneous windows when the profile transitions from one specific state to another.

- **Arguments:**

- `from` : instance of `Schema` representing the state the profile must transition from
- `to` : instance of `Schema` representing the state the profile must transition from

- **Returns:** `Windows`

method `changes`

Computes the instantaneous windows when the profile changes

- **Returns:** `Windows`

method `equal`

Computes the windows when this profile is equal to another

- **Arguments:**
 - `other` : `Discrete<Schema>`
- **Returns:** `Windows`

method `notEqual`

Computes the windows when this profile is not equal to another

- **Arguments:**
 - `other` : `Discrete<Schema>`
- **Returns:** `Windows`

Precomputed Resource Profiles

The What

Precomputed profiles are quite similar to the profiles produced by a Merlin simulation. They describe the behavior of some dynamic characteristic over time and can be viewed in the Aerie UI along with a plan. The big difference between precomputed profiles as those that come from simulation is that precomputed profiles are... well... precomputed. Precomputed profiles come from the user, rather than Merlin's simulation capabilities, and can be uploaded at any time, before or after simulation.

The Why

Ultimately, the purpose of precomputed profiles is up to the user. You might be looking to include some geometry information to assist in building a plan, or you may simply be adding power and thermal modeling results to be viewable with existing simulation results. Just be aware that at this time precomputed profiles are not accessible to the mission model during simulation, and are simply for viewing purposes in the UI.

The How

To upload a set of precomputed profiles, you will need to add what is called an external dataset to a plan. An external dataset is just a set of profiles with a start time from which the profiles are based. The dataset must be associated with a plan, and will persist as long as the plan exists, or until it is deleted by a user. An external dataset can be added via the GraphQL mutation `addExternalDataset`.

GraphQL Mutation: `addExternalDataset`

The `addExternalDataset` GraphQL mutation takes three parameters as specified below:

- `planId` : Type: Int Description: The ID of the plan to associate the external dataset with.
- `datasetStart` : Type: String Description: The UTC timestamp the dataset is based from. UTC Format: `yyyy-dddThh:mm:ss`
- `profileSet` : Type: Object Description: The set of precomputed profiles that make up the external dataset.

The profile set to be uploaded should have one entry for each precomputed profile, indexed by a unique name mapping to an object specifying the details of the profile. Each profile should have a `type` field, which specifies whether the profile is real- or discrete-valued. A discrete profile must also contain a `schema` field, which specifies the schema of the values it takes on. These are not limited to basic types, but can take on any complex structure made up using our `ValueSchema` construct (for more information, see our [ValueSchema documentation](#)).

Finally, both profile types require a list of segments that describe the actual behavior of the profile. The `segments` field of each profile maps to a list of objects, each of which details a segment. Each segment should contain the following two fields:

- `duration` : Type: Integer Description: The duration (in microseconds) the segment's dynamics hold before the next segment begins.
- `dynamics` : Type: Dependent on profile type. Description: The behavior of the profile over the lifetime of this segment.

A discrete profile's dynamics should match the format specified by the `schema` field, while a real profile's dynamics should always contain an initial value and a rate of change. See our example external dataset mutations [below](#) to see both profile specification types.

Deleting External Datasets

There may be a time when you find an external dataset you've been using is no longer relevant, and must be removed. This is easily done by providing the ID of the dataset you wish to delete to the `delete_dataset_by_pk` mutation. For example, the following mutation will delete the dataset with id 5:

```
mutation deletedDataset {
  delete_dataset_by_pk(id: 5) {
    id
  }
}
```

Example External Dataset Mutations

Example 1

This example shows an external dataset being uploaded to the plan with ID `2` starting at `2018-331T04:00:00`. Two precomputed profiles are

included in the external dataset. First a real profile called `batteryEnergy` starts at a value of 50 and decreases at a rate of -0.5 units per second over 30 seconds. At that point, the value is 35 and the rate is changed to -0.1 units per second for 30 more seconds. The second profile, a discrete profile called `awake` contains a schema that tells us its values are boolean. The segments tell us that for the first 30 seconds the profile's dynamics are the value `true` and for the next 30 seconds the value `false`. At the end we see the ID of the external dataset that is created is queried as the result of this mutation.

```
mutation {
  addExternalDataset(
    planId: 2
    datasetStart: "2018-331T04:00:00"
    profileSet: {
      batteryEnergy: {
        type: "real"
        segments: [
          { duration: 30000000, dynamics: { initial: 50, rate: -0.5 } }
          { duration: 30000000, dynamics: { initial: 35, rate: -0.1 } }
        ]
      }
      awake: {
        type: "discrete"
        schema: { type: "boolean" }
        segments: [
          { duration: 30000000, dynamics: true }
          { duration: 30000000, dynamics: false }
        ]
      }
    }
  ) {
    datasetId
  }
}
```

Example 2

This mutation adds to plan with id `7` an external dataset starting at `2038-192T14:00:00` with a single precomputed profile, `orientation`. This discrete profile's schema tells us that its values are structs with real-valued `x`, `y` and `z` fields. For the first hour, the profile takes a value of `x=0`, `y=0`, `z=1`. For the next hour thereafter, the profile is valued at `x=1`, `y=1`, `z=0`.

```
mutation {
  addExternalDataset(
    planId: 7
    datasetStart: "2038-192T14:00:00"
    profileSet: {
      orientation: {
        type: "discrete"
        schema: {
          type: "struct"
          items: {
            x: { type: "real" }
            y: { type: "real" }
            z: { type: "real" }
          }
        }
      }
      segments: [
        { duration: 3600000000, dynamics: { x: 0, y: 0, z: 1 } }
        { duration: 3600000000, dynamics: { x: 1, y: 1, z: 0 } }
      ]
    }
  ) {
    datasetId
  }
}
```


Scheduling

Note: Automated scheduling eDSL is under heavy development. This documentation will change regularly.

- [Introduction](#)
 - [Authoring a new goal](#)
 - [Specifying the order of goals](#)
 - [Running the scheduler](#)
- [Scheduling DSL Documentation](#)
 - [Activity Template](#)
 - [Goal Types](#)
 - [Activity Recurrence Goal](#)
 - [Coexistence Goal](#)
 - [Cardinality Goal](#)
 - [Restricting when a goal is applied](#)

Introduction

This guide explains how to use the scheduling service with the latest version of Aerie.

The scheduling service allows you to add activities to a plan based on goals that you define. These goals are defined in [TypeScript](#), using a library provided by Aerie.

This guide assumes some very basic familiarity with programming (terminology like functions, arguments, return values, types, and objects will not be explained), but it does not assume the reader has seen TypeScript or JavaScript before.

Prerequisites

- You will need to have uploaded a Mission Model
- You will need to have created a plan from that mission model

Steps

1. [Authoring a new goal](#)
2. [Specifying the order of goals](#)
3. [Running the scheduler](#)

Authoring a new goal

In the Aerie UI, open the scheduling pane, and click [New Goal](#). This will open a text editor, with the following default text:

```
export default (): Goal => {  
  // Your code here  
}
```

This is a TypeScript function that takes no arguments and returns a Goal.

To unpack all of the parts:

- `export default` signals to Aerie that this is the function that defines the Goal.
- `() => {}` in TypeScript is called an [arrow function](#).
 - The parenthesis `()` represent the parameters that the function takes. Scheduling goals cannot take any parameters, so these parenthesis must be empty.

- The code provided when you click **New Goal** is incomplete - the function does not yet return a Goal, so you should see the word **Goal** underlined

Mousing over the word Goal, you should see something akin to the following message:

This message means that the function has promised to return a value, but it currently lacks a return statement. Between the curly braces, add the following code: `return Goal.ActivityRecurrenceGoal()` :

Now, the editor should tell you that `ActivityRecurrenceGoal()` takes one argument.

The argument that we're missing is the "options" object. Objects in typescript are defined using curly braces `{ }` with key-value pairs, like so:

If we pass an empty object `{}` to `ActivityRecurrenceGoal`, we will get a new error message that will tell us what keys our object will need:

Argument of type '{}' is not assignable to parameter of type '{ activityTemplate: ActivityTemplate; interval: number; }'.

Type '{}' is missing the following properties from type '{ activityTemplate: ActivityTemplate; interval: number; }': activityTemplate, [interval](#)

This error message tells us that our object is missing two keys: `activityTemplate`, and `interval`. If we look up the definition of `ActivityRecurrenceGoal` in the scheduling documentation, we will see that it does indeed need an activity template and an interval. Let's add those:

Now, we just need to finish specifying the **activityTemplate**. Start by typing `ActivityTemplates.` (note the period), and select an activity type.

Provide your activity an object with the arguments that that activity takes. Once the editor is no longer underlining your code, save your goal, and hit `Schedule & Analyze` and observe that your new activities were added to the plan.

Specifying the order of goals

The Aerie scheduler accepts a list of goals, and tries to satisfy them one by one by adding activities to your plan. We refer to this list of goals as a **scheduling specification**. Currently, Aerie creates one scheduling specification per plan. A goal's **priority** is simply a number reflecting that goal's position in the scheduling specification. The first goal will always have priority `0`, and the n-th goal will always have priority `n - 1`. Within a scheduling specification, a scheduling goal can be toggled enabled or disabled. A disabled scheduling goal will be excluded from scheduling execution. The priority ordering of the scheduling goals remains unchanged whether a goal is enabled or disabled.

Scheduling DSL Documentation

Here you will find the full set of features in the scheduling DSL.

Known General Issues

Activities with `uncontrollable durations` have been found to behave somewhat unpredictably, in terms of when they are placed. This has to do with how temporal constraints interact with the unpredictability of the durations. Finding when an activity will start while subject to temporal constraint involves search.

ActivityTemplate

An `ActivityTemplate` specifies the type of an activity, as well as the arguments it should be given. Activity templates are generated for each mission model. You can get the full list of activity templates by typing `ActivityTemplates.` (note the period) into the scheduling goal editor, and viewing the auto-complete options.

If the activity has parameters, pass them into the constructor in a dictionary as key-value pairs (i.e. `ActivityTemplate.ParamActivity({ param:1 })`). If the activity has no parameters, do not pass a dictionary (i.e. `ActivityTemplate.ParameterlessActivity()`).

Goal Types

Activity Recurrence Goal

The Activity Recurrence Goal (sometimes referred to as a "frequency goal") specifies that a certain activity should occur repeatedly throughout the plan, at some given interval.

Inputs

- **activityTemplate**: the description of the activity whose recurrence we're interested in.
- **interval**: a `Duration` of time specifying how often this activity must occur

Behavior

This interval is treated as an upper bound - so if the activity occurs more frequently, that is not considered a failure.

The scheduler will find places in the plan where the given activity has not occurred within the given interval, and it will place an instance of that activity there.

Note: The interval is measured between the *start times* of two activity instances. Neither the duration, nor the end time of the activity are examined by this goal.

Example

```
export default function myGoal() {
  return Goal.ActivityRecurrenceGoal({
    activityTemplate: ActivityTemplates.GrowBanana({
      quantity: 1,
      growingDuration: 1 * 60 * 60 * 1000 * 1000, //1 hour in microseconds
    }),
    interval: 2 * 60 * 60 * 1000 * 1000 // 2 hours in microseconds
  })
}
```

The goal above will place a `GrowBanana` activity in every 2-hour period of time that does not already contain one with the exact same parameters.

Coexistence Goal

The Coexistence Goal specifies that a certain activity should occur once **for each** occurrence of some condition.

Inputs

- **forEach**: a set of time windows (`Windows` , see [documentation](#) on how to produce such an expression) or a set of activities (`ActivityExpression`)
- **activityTemplate**: the description of the activity to insert after each activity identified by `forEach`
- **startsAt**: optionally specify a specific time when the activity should start relative to the window
- **startsWithin**: optionally specify a range when the activity should start relative to the window
- **endsAt**: optionally specify a specific time when the activity should end relative to the window
- **endsWithin**: optionally specify a range when the activity should end relative to the window

NOTE: Either the start or end of the activity must be constrained. This means that at least **1** of the 4 properties `startsAt` , `startsWithin` , `endsAt` , `endsWithin` must be given.

Behavior

The scheduler will find places in the plan where the `forEach` condition is true, and if not, it will insert a new instance using the given `activityTemplate` and temporal constraints.

Examples

```
export default () => Goal.CoexistenceGoal({
  forEach: ActivityExpression.ofType(ActivityTypes.GrowBanana),
  activityTemplate: ActivityTemplates.PeelBanana({peelDirection: "fromStem"}),
  startsAt: TimingConstraint.singleton(WindowProperty.END).plus(5 * 60 * 1000 * 1000)
})
```

Behavior: for each activity A of type `GrowBanana` present in the plan when the goal is evaluated, place an activity of type `PeelBanana` starting exactly at the end of A + 5 minutes.

```
export default () => Goal.CoexistenceGoal({
  forEach: ActivityExpression.ofType(ActivityTypes.GrowBanana),
  activityTemplate: ActivityTemplates.PeelBanana({peelDirection: "fromStem"}),
  startsWithin: TimingConstraint.range(WindowProperty.END, Operator.PLUS, 5 * 60 * 1000 * 1000),
  endsWithin: TimingConstraint.range(WindowProperty.END, Operator.PLUS, 6 * 60 * 1000 * 1000)
})
```

Behavior: for each activity A of type `GrowBanana` present in the plan when the goal is evaluated, place an activity of type `PeelBanana` starting in the interval [end of A, end of A + 5 minutes] and ending in the interval [end of A, end of A + 6 minutes].

The following diagram shows an example of how the temporal constraints defined in this example affects the placement of a `PeelBanana` activity with respect to an existing `GrowBanana` activity.

```
gantt
  dateFormat HH:mm
  axisFormat %H:%M
  GrowBanana : 17:45, 10min
  Start interval for PeelBanana : 17:55, 5min
  End interval for PeelBanana : 17:55, 8min
  Possible instantiation of PeelBanana : 17:56, 2min
  Other possible instantiation of PeelBanana : 17:59, 3min
```

```
export default () => Goal.CoexistenceGoal({
  forEach: Real.Resource("/fruit").equal(4.0),
  activityTemplate: ActivityTemplates.PeelBanana({peelDirection: "fromStem"}),
  endsAt: TimingConstraint.singleton(WindowProperty.END).plus(5 * 60 * 1000 * 1000)
})
```

Behavior: for each continuous period of time during which the `/fruit` resource is equal to 4, place an activity of type `PeelBanana` ending exactly at the end of A + 6 minutes. Note that the scheduler will allow a default timing error of 500 milliseconds for temporal constraints. This parameter will be configurable in an upcoming release.

KNOWN ISSUE: If the end is unconstrained while the activity has an uncontrollable duration, the scheduler may fail to place the activity. To work

around this, add an `endsWithin` constraint that encompasses your expectation for the duration of the activity - this will help the scheduler narrow the search space.

Cardinality Goal

The Cardinality Goal specifies that a certain activity should occur in the plan either a certain number of times, or for a certain total duration.

Inputs

- **activityTemplate**: the description of the activity whose recurrence we're interested in.
- **specification**: an object with either an `occurrence` field, a `duration` field, or both (see examples below).

Behavior

The duration and occurrence are treated as lower bounds - so if the activity occurs more times, or for a longer duration, that is not considered a failure, and the scheduler will not add any more activities.

The scheduler will identify whether it not the plan has enough occurrences or total duration of the given activity template. If not, it will add activities until satisfaction.

Examples

Setting a lower bound on the total duration:

```
export default function myGoal() {
  return Goal.CardinalityGoal({
    activityTemplate: ActivityTemplates.GrowBanana({
      quantity: 1,
      growingDuration: 1000000,
    }),
    specification: { duration: 10 * 1000000 }
  })
}
```

Setting a lower bound on the number of occurrences:

```
export default function myGoal() {
  return Goal.CardinalityGoal({
    activityTemplate: ActivityTemplates.GrowBanana({
      quantity: 1,
      growingDuration: 1000000,
    }),
    specification: { occurrence: 10 }
  })
}
```

Combining the two:

```
export default function myGoal() {
  return Goal.CardinalityGoal({
    activityTemplate: ActivityTemplates.GrowBanana({
      quantity: 1,
      growingDuration: 1000000,
    }),
    specification: { occurrence: 10, duration: 10 * 1000000 }
  })
}
```

NOTE: In order to avoid placing multiple activities at the same start time, the Cardinality goal introduces an assumed mutual exclusion constraint - namely that new activities will not be allowed to overlap with existing activities.

OR goal - Disjunction of goals

The OR Goal aggregates several goals together and specifies that at least one of them must be satisfied.

Inputs

- **goals**: a list of goals (here below referenced as the subgoals)

Behavior

The scheduler will try to satisfy each subgoal in the list until one is satisfied. If a subgoal is only partially satisfied, the scheduler will not backtrack and will let the inserted activities in the plan.

Examples

```
export default function myGoal() {
  return Goal.CardinalityGoal({
    activityTemplate: ActivityTemplates.GrowBanana({
      quantity: 1,
      growingDuration: 1000 * 1000 * 60 * 60, //1 hour in microseconds
    }),
    specification: { occurrence : 10 }
  }).or(
    Goal.ActivityRecurrenceGoal({
      activityTemplate: ActivityTemplates.GrowBanana({
        quantity: 1,
        growingDuration: 1 * 60 * 60 * 1000 * 1000, //1 hour in microseconds
      }),
      interval: 2 * 60 * 60 * 1000 * 1000 // 2 hours in microseconds
    })
  ))
}
```

If the plan has a 24-hour planning horizon, the OR goal above will try placing activities of the `GrowBanana` type. The first subgoal will try placing 10 1-hour occurrences. If it fails to do so, because the planning horizon is maybe too short, it will then try to schedule 1 activity every 2 hours for the duration of the planning horizon.

It may fail to achieve both subgoals but as the scheduler does not backtrack for now, activities inserted by any of the subgoals are kept in the plan.

AND goal - Conjunction of goals

The AND Goal aggregates several goals together and specifies that at least one of them must be satisfied.

Inputs

- **goals**: an ordered list of goals (here below referenced as the subgoals)

Behavior

The scheduler will try to satisfy each subgoal in the list. If a subgoal is only partially satisfied, the scheduler will not backtrack and will let the inserted activities in the plan. If all the subgoals are satisfied, the AND goal will appear satisfied. If one or several subgoals have not been satisfied, the AND goal will appear unsatisfied.

Examples

```
export default function myGoal() {
  return Goal.CoexistenceGoal({
    forEach: Real.Resource("/fruit").equal(4.0),
    activityTemplate: ActivityTemplates.PeelBanana({peelDirection: "fromStem"}),
    endsAt: TimingConstraint.singleton(WindowProperty.END).plus(5 * 60 * 1000 * 1000)
  }).and(
    Goal.CardinalityGoal({
      activityTemplate: ActivityTemplates.PeelBanana({peelDirection: "fromStem"}),
      specification: { occurrence : 10 }
    })
  ))
}
```

The AND goal above has two subgoals. The coexistence goal will place activities of type `PeelBanana` everytime the `/fruit` resource is equal to 4. The second goal will place 10 occurrences of the same kind of activities `PeelBanana`. The first subgoal will be evaluated first and will place a certain number of `PeelBanana` activities in the plan. When the second goal will be evaluated, it will count already present `PeelBanana` activities and insert the missing number. Imagine the first goals leads to inserting 2 activities. The second goal will then have to place 8 activities to be satisfied.

Restricting when a goal applied

By default, a goal applies on the whole planning horizon. The Aerie scheduler provides support for restricting *when* a goal applies with the `.applyWhen()` method in the `Goal` class. This node allows users to provide a set of windows (`Windows` , see [documentation](#)) which could be a time or a resource-based window.

The `.applyWhen()` method, takes one argument: the windows (in the form of an expression) that the goal should apply over. What follows is an example that applies a daily recurrence goal only when a given resource is greater than 2. If the resource is less than two, then the goal is no

longer applied.

```
export default function myGoal() {
  return Goal.ActivityRecurrenceGoal({
    activityTemplate: ActivityTemplates.GrowBanana({
      quantity: 1,
      growingDuration: 1 * 60 * 60 * 1000 * 1000, //1 hour in microseconds
    }),
    interval: 2 * 60 * 60 * 1000 * 1000 // 2 hours in microseconds
  }).applyWhen(Real.Resource("/fruit").greaterThan(2))
}
```

Notes on boundaries:

- If you are trying to schedule an activity, or a recurrence within a window but that window cuts off either the activity or the recurrence interval (depending on the goal type), it will not be scheduled. For example, if you had a recurrence interval of 3 seconds, scheduling a 2 second activity each recurrence, and had the following window, you'd get the following:

```
RECURRENCE INTERVAL: [+++++---]
GOAL WINDOW:         [+++++---]
RESULT:              [+++++---]
```

That, is, the second activity won't be scheduled as the goal window cuts off its recurrence interval.

- Scheduling is *local*, not global. This means for every window that is matched (as it is possible to have disjoint windows, imagine a resource that fluctuates upward and downward but only applying that goal when the resource is over a certain value), the goal is applied individually. So, for that same recurrence interval setup as before, we could have:

```
RECURRENCE INTERVAL: [+++++---]
GOAL WINDOW:         [+++++---]
RESULT:              [+++++---] //(the second one is applied independently of the first!)
```

- When mapping out a temporal window to apply a goal over, keep in mind that the ending boundary of the goal is *exclusive*, i.e. if I want to apply a goal in the window of 10-12 seconds, it will apply only on seconds 10 and 11. This is in line with the [fencepost problem](#).

Command Expansion Guide

Introduction

This guide explains how to use the command expansion service with the latest version of Aerie.

The command expansion service allows you to generate commands from activity information within an Aerie Plan. With this service, you can create custom logic that can be assigned to an activity_type. This custom logic will be used by Aerie to generate the commands. Your custom logic will require some pre-generated libraries provided by the command expansion service. An IDE will be used to write your custom logic.

Prerequisites

- An AMPCS Command Dictionary
- Mission Model

Steps

1. Upload a Command Dictionary and retrieve the command_library.ts file
2. Select an activity_type to generate the activity_library.ts file
3. Setup an IDE to start writing custom logic
4. Upload the custom logic
5. Create an Expansion Set
6. Expand the plan with the Expansion Set
7. Retrieve the generated commands

Continue to [Upload and Generate Command Typescript Library](#)

Upload and Generate Command Typescript Library

We recommend using the latest command dictionary as a starting point. This will allow you to use the most up-to-date commands that have been tested and validated. When you upload a command dictionary, an autogenerated `command_library` typescript file will be available to download. This library will be used later.

Upload

Uploading a command dictionary is very simple using GraphQL. In order to upload a command dictionary, login to the Aerie UI and navigate to the GraphQL Playground. From there, you will use our pre-defined mutation to upload the command dictionary. Run this GraphQL mutation below to upload the command dictionary.

```
mutation Upload(  
  $commandDictionaryString: String!  
) {  
  uploadDictionary(dictionary: $commandDictionaryString) {  
    id  
    command_types_typescript_path  
    mission  
    version  
    created_at  
    updated_at  
  }  
}
```

Return Values

- id - unique id of the database entry
- command_types_typescript_path - path where the generated command library is stored for the upload dictionary
- mission - the mission name for the command dictionary uploaded
- version - The version of the command dictionary uploaded
- created_at - when was the command dictionary first uploaded
- updated_at - when was the command dictionary updated thus generating a new command library

Download Command Library

A successful upload will return an `id` used to retrieve the library. Running the GraphQL query below will retrieve the typescript library.

1. Run this GraphQL command below

```
query GetCommandTypescript($commandDictionaryId: String!) {  
  getCommandTypeScript(commandDictionaryId: commandDictionaryId) {  
    reason  
    status  
    typescriptFiles {  
      content  
      filePath  
    }  
  }  
}
```

2. Save this to a file named `command-types.ts`

Return Values

- content - The typescript library for your activity_type
- filepath - the filename of your library
- reason - If any errors were thrown when generating this library file, the error message would be stored here. If there are no errors, you will get a `null`
- status - server-status ex. 404, 500, 200

Continue to [Generate Activity Typescript Library](#)

Generate Activity Typescript Library

In order to write an author expansion logic, you will need to know which `activity_types` to use. These `activity_types` are defined in the mission model.

Upon the selection of an `activity_type`, the command expansion service will generate a typescript library file containing all parameters that can be used when developing your logic. Below is the GraphQL request used to retrieve the library.

1. Run the graphQL query below

```
query GetActivityTypescript($activityTypeName: String!, $missionModelId: Int!) {  
  getActivityTypeScript(activityTypeName: $activityTypeName, missionModelId: $missionModelId) {  
    typescriptFiles {  
      content  
      filePath  
    }  
    reason  
    status  
  }  
}
```

2. Save to a file <activity_type>_activity-types.ts ex. BakeBananaBread_activity-types.ts

Return Values

- content - The typescript library for your activity_type
- filepath - the filename of your library
- reason - If any errors were thrown when generating this library file, the error message would be stored here. If there are no errors, you will get a `null`
- status - server-status ex. 404, 500, 200

Continue to [Setup IDE](#)

IDE Setup

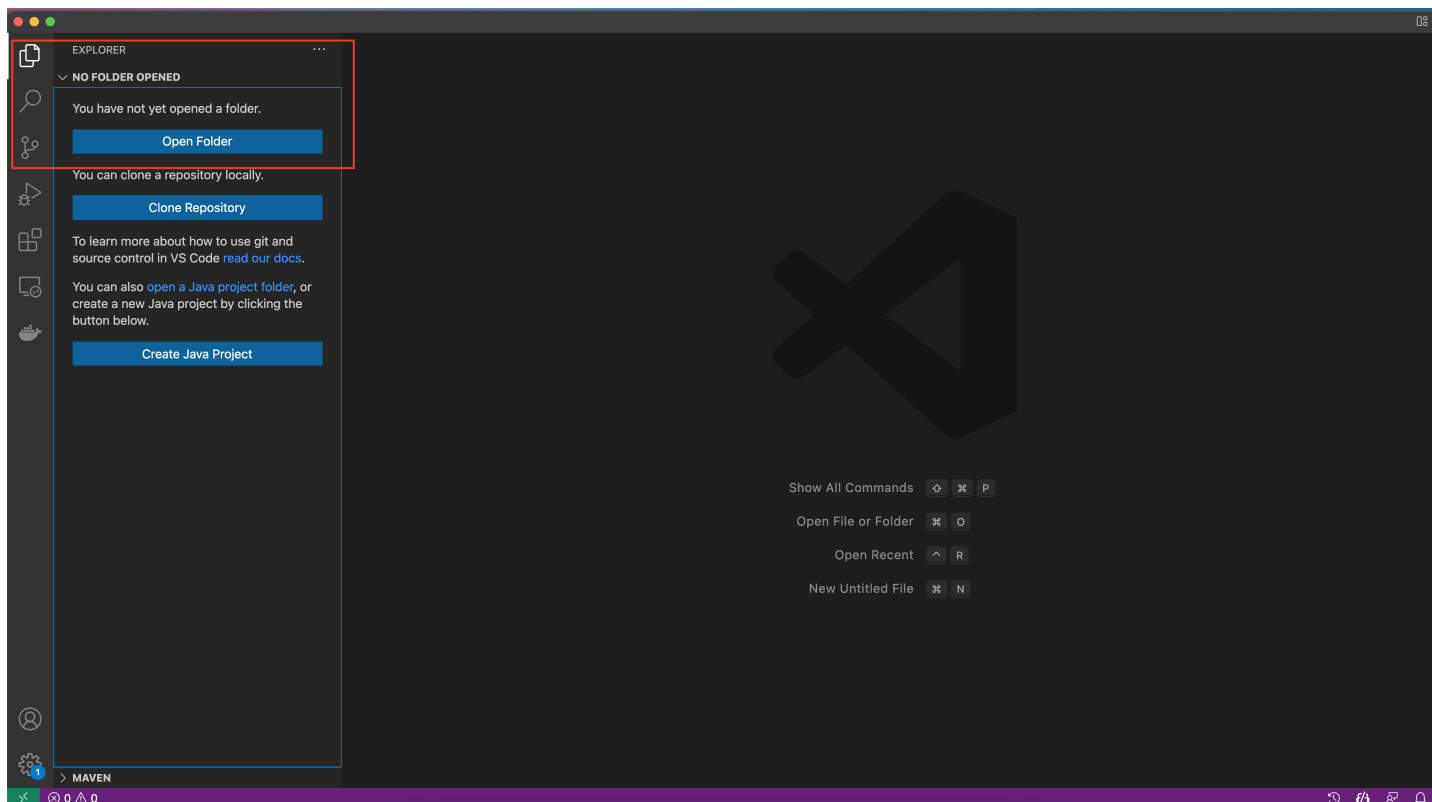
Now it is time to start writing the expansion logic for the activity type. The recommended IDE is Visual Studio Code as we are looking at embedding this IDE as a web editor as a future service. Another option is IntelliJ but you will need the paid Ultimate version to use the typescript plugin.

Recommended IDE:

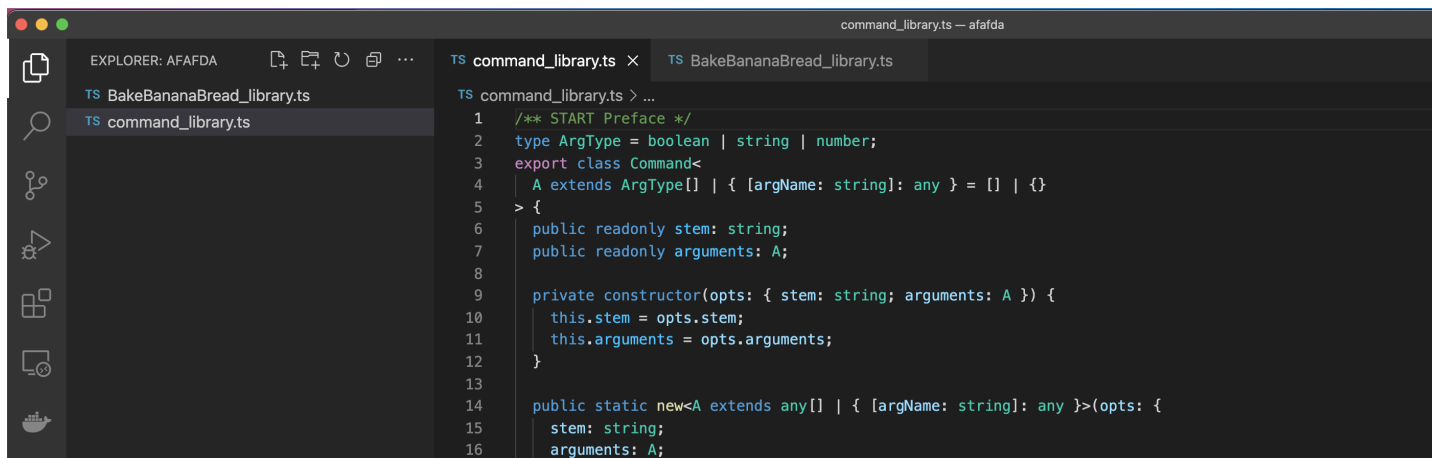
- Visual Studio Code: <https://code.visualstudio.com/>
- IntelliJ: <https://www.jetbrains.com/idea/>

Setup:

1. Open a new project folder. You can create a new folder as your workspace



2. Drag you <activity_library>.ts and <command_library>.ts into you project



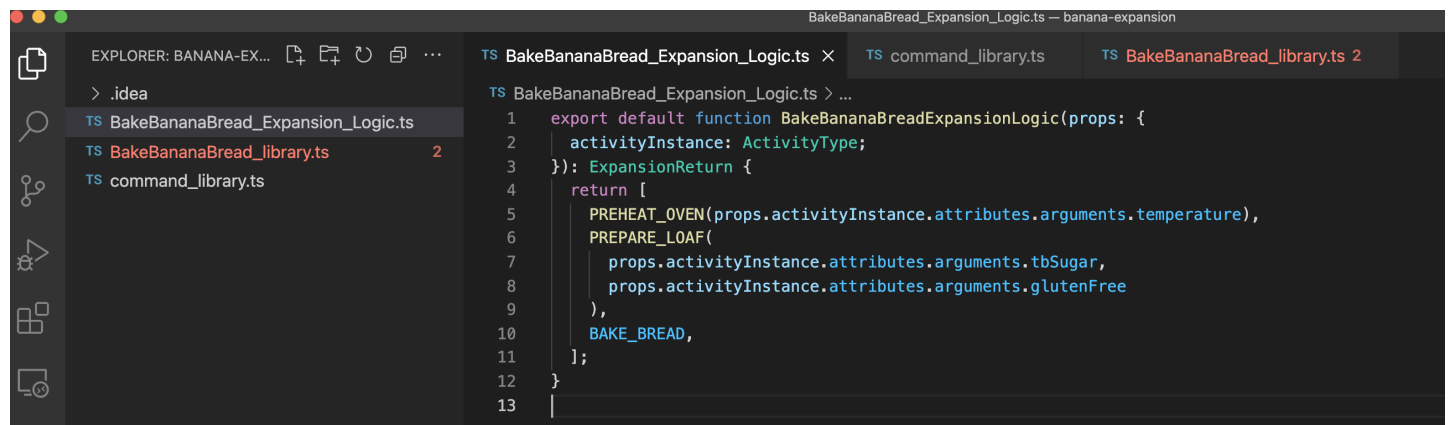
3. Create an expansion logic file. You can name it whatever you like. **ex. BakeBananaBreadExpansionLogic.ts**
4. Add the following below which is a boilerplate template

```
// This is the entry point for the expansion - name it whatever you want, it just has to be the default export
export default function CommandExpansion(
  // Everything related to this activity
  props: {
    activityInstance: ActivityType;
  }
): ExpansionReturn {
  // Commands are fully strongly typed, and intellisense in the authoring editor will
  // guide users to write correct expansions
  return [
    // COMMAND_A(arg1)
    Subroutine(),
    // Commands without any arguments omit the ()
    // COMMAND_C
  ];
}

// You can break it into smaller logical/reusable chunks to make comprehension easier
function Subroutine() {
  return [
    //COMMAND_B(arg1,arg2)
  ];
}
```

Writing

As you start typing in the editor, notice that each command is validated and type-checked. This means as you write out your logic, the editor will ensure all commands are valid, and that each argument meets the requirements. When you are happy with your logic save your works.



Continue to [Expansion Logic API](#)

Expansion Logic API

Expansion Logic API

Here is a list of APIs that AERIE provides to help you build your expansion logic.

Activity's Attributes

You have access to the activity's attributes such as `start time`, `end time`, `duration`, etc

```
C.ECHO(`Activity's Start Time: ${props.activityInstance.startTime}`), // YYYY-DDDThh:mm:ss
C.ECHO(`Activity's End Time: ${props.activityInstance.endTime}`), // YYYY-DDDThh:mm:ss
C.ECHO(`Activity's Start Offset: ${props.activityInstance.startOffset}`), // hh:mm:ss
C.ECHO(`Activity's Duration: ${props.activityInstance.Duration}`), // hh:mm:ss
```

Computed Attributes

You have access to the activity type's computed attributes. You can use these computed attributes in your expansion logic

```
export default function CommandExpansion(props: {
  activityInstance: ActivityType;
}): ExpansionReturn {
  return [
    PeelBanana(),
    C.PREHEAT_OVEN(350),
    C.PREPARE_LOAFF(1, false),
    C.BAKE_BREAD,
  ];

  function PeelBanana(): Command {
    if (props.activityInstance.attributes.computed < 2) {
      return C.PEEL_BANANA("fromStem");
    } else {
      return C.ECHO("Already have enough Banana's peeled...");
    }
  }
}
```

Time

You can specify a time argument for your commands. The time arguments supported are `absolute`, `relative`, and `epoch`. You will be using the Javascript Temporal Polyfill to represent time values. <https://tc39.es/proposal-temporal/docs/>

- Absolute Time (`Temporal.Instant`): YYYY-DDDThh:mm:ss
- Relative Time (`Temporal.Duration`): hh:mm:ss
- Epoch Time (`Temporal.Duration`): hh:mm:ss

```
// Absolute examples
A`2020-060T03:45:19`.ADD_WATER
A(Temporal.Instant.from("2025-12-24T12:01:59Z")).ADD_WATER
A("2020-060T03:45:19").ADD_WATER

// Relative examples
R`00:15:00`.EAT_BANANA
R(Temporal.Duration.from({ minutes: 15 })).EAT_BANANA
R('00:15:00').EAT_BANANA

// Epoch examples
E`00:15:30`.PREPARE_LOAF(1,true)
E(Temporal.Duration.from({ minutes: 15, seconds: 30 })).PREPARE_LOAF(1,true)
E('00:15:30').PREPARE_LOAF(1,true)

// Command Complete examples
C.BAKE_BREAD
```

Submit

Continue to [Submit Expansion Logic](#)

Submit Expansion Logic

Upload

Once you've defined your expansion logic you need to upload it back to the commanding server and link it to the `activity_type` you wrote it for. Below you can use the GraphQL to upload the expansion logic. Use the GraphQL to upload the expansion logic file for the activity_type

```
mutation UploadExpansionLogic(
  $activityTypeName: String!
  $expansionLogic: String!
) {
  addCommandExpansionTypeScript(
    activityTypeName: $activityTypeName
    expansionLogic: $expansionLogic
  ) {
    id
  }
}
```

Retrieve

If you would like to retrieve an expansion logic for modification use this GraphQL API below:

```
query GetExpansionLogic(
  $expansionRuleId: Int!
) {
  expansion_rule(where: {id: {_eq: $expansionRuleId}}) {
    expansion_logic
    activity_type
    id
  }
}
```

Retrieve a specific activity_type's various expansion logics:

```
query GetExpansionLogic(
  $activityType: String!
) {
  expansion_rule(where: {activity_type: {_eq: $activityType}}) {
    expansion_logic
    activity_type
    id
  }
}
```

Continue to [Create an Expansion Set](#)

Create an Expansion Set

Introduction

Think of an "expansion_set" as a way of grouping expansion logic to different versions of mission models or command dictionaries.

But why the need for an "expansion_set"?

For example, say you have a mission model and command dictionary version 1. You have written expansion logic that uses commands from the command dictionary. Later the command dictionary or mission model is updated to version 2. If your original expansion logic was created to only be compatible with the previous version of the command dictionary or mission model, you need to go through and rewrite/update your existing logic for all the activity_types.

An alternative is that you can use the "expansion_set". This will allow you to copy over a set of expansion logic into a new "expansion_set" without manually changing anything in your code. Some expansion logic might not be valid anymore but the user can decide which expansion logic to copy over to the new set containing the updated command dictionary or mission model.

A future feature will inform users of which expansion rules are compatible/incompatible for the selected mission model and command dictionary, making the migration of expansion_sets easier.

The below GraphQL creates an expansion set for a given command dictionary and mission model which includes the provided expansions.

```
mutation CreateExpansionSet(  
  $commandDictionaryId: Int!  
  $missionModelId: Int!  
  $expansionIds: [Int!]!  
) {  
  createExpansionSet(  
    commandDictionaryId: $commandDictionaryId  
    missionModelId: $missionModelId  
    expansionIds: $expansionIds  
  ) {  
    id  
  }  
}
```

Ex. 1

That was a lot to take in. Let's create our own expansion set and tie everything together. Below is an example GraphQL to create an "expansion_set". With the 4 expansion logic defined we want to group them with command dictionary and mission model version 1.

```
{  
  "commandDictionaryId": 1,  
  "missionModelId": 1,  
  "expansionIds": [1,2,4,9]  
}
```

Ex.2

Now let us say the mission model has changed and out of the 4 activity_types defined in the mission model we are dropping one of them. We will only copy over the expansion logic that is valid in this new set.

```
{  
  "commandDictionaryId": 1,  
  "missionModelId": 2,  
  "expansionIds": [1,2,9]  
}
```

List Expansion Sets

Below is a way to list the created expansion sets in GraphQL:


```

query GetAllExpansionSets {
  expansion_set {
    id
    command_dictionary {
      version
    }
    mission_model {
      id
    }
    expansion_rules {
      activity_type
      id
    }
  }
}

```

Get all `expansion_set`'s for a particular `mission_model` or `command dictionary`:

```

query GetSpecificExpansionSet(
  $commandDictionaryId: Int!
  $missionModelId: Int!
) {
  expansion_set(where: {command_dictionary: {id: {_eq: $commandDictionaryId}}, mission_model_id: {_eq: $missionModelId}}) {
    id
    expansion_rules {
      activity_type
      id
    }
  }
}

```

Continue to [Expand the Plan](#)

Expand the Plan

Introduction

Now for the exciting part! We can finally expand the plan into commands. To do that, we'll need to set up a couple of things with your plan.

1. Add an `activity_instances` to the plan which has expansion logic defined in the `expansion_set`
2. Run a simulation on the plan

Prerequisite

At this time you will have to use the Hasura Action API in order to expand the plan into commands. You will have to retrieve some information needed to expand via the API. The following two values are needed:

- `expansionSetID`
- `simulationDatasetId`

Below is the query to retrieve an `expansionSetID`:

```
query GetExpansionSet(
  $commandDictionaryId: Int!
  $missionModelId: Int!
) {
  expansion_set(where: { mission_model_id: { _eq: $missionModelId }, command_dictionary: { id: { _eq: $commandDictionaryId } } }) {
    id
  }
}
```

id : 1

Below is the query to retrieve the `simulationDatasetId`:

```
query GetSimulationDatasetId(
  $planId: Int!
) {
  simulation(where: { plan_id: { _eq: $planId } }, order_by: { dataset: { id: desc } }, limit: 1) {
    dataset {
      id
    }
  }
}
```

id : 5

Expanding

Below is an example of the Hasura Action you can use to expand a plan:

```
mutation ExpandPlan(
  $expansionSetId: Int!
  $simulationDatasetId: Int!
) {
  expandAllActivities(expansionSetId: $expansionSetId, simulationDatasetId: $simulationDatasetId) {
    id
    expandedActivityInstances {
      commands {
        stem
      }
      errors {
        message
      }
    }
  }
}
```

Continue to [Get Expanded Commands](#)

Get Expanded Commands

When the plan has been expanded you will be able to retrieve the generated commands from the expansion run. Each `activity_instance` will have commands and any errors generated. If there are no commands and errors listed then an `activity_instance` didn't have any expansion logic defined. Below is an example of how to retrieve all the expanded commands from an expansion run or a specific command from `activity_instance`.

```
query GetAllExpandedCommandsForExpansionRun(
  $expansionRunId: Int!
) {
  activity_instance_commands(where: {expansion_run: {id: {_eq: $expansionRunId}}}) {
    activity_instance_id
    commands
    errors
  }
}
```

```
query GetCommandsForExpansionRunAndActivityInstance(
  $expansionRunId: Int!
  $activityInstanceId: Int!
) {
  activity_instance_commands(where: {activity_instance_id: {_eq: $activityInstanceId}, expansion_run: {id: {_eq: $expansionRunId}}})
  {
    commands
    errors
  }
}
```

Continue to [Create Sequence](#)

Create Sequence

Sequence creations starts with declaring a new sequence on a simulation dataset. Because of the complexities with activities in the plans regenerating on simulation, sequences are tied to a specific simulation run rather than a plan (future work will explore addressing this).

The API call to do this is a simple mutation

```
mutation CreateSequence(  
  $seqId: String!  
  $simulationDatasetId: Int!  
) {  
  insert_sequence_one(object: {  
    simulation_dataset_id: $simulationDatasetId,  
    seq_id: $seqId,  
    metadata: {},  
  }) {  
    seq_id  
  }  
}
```

Where the seqId is the id you want assigned to the sequence and simulationDatasetId is the id of the simulation dataset for the simulation you wish to create the sequence for.

Note: You can only have one of any given seqId for a given simulationDataset

Continue to [Link Simulated Activity to Sequence](#)

Link Simulated Activity to Sequence

Once a sequence is created, you can link simulated activities to the sequence. Linking simulated activities to a sequence specifies that that activity's commands should be included in the sequence. Note that a simulated activity is NOT the same as an activity inserted in the plan, but the result of simulation on that plan and the ids are NOT interchangeable. You can retrieve the simulated activities for a simulation dataset to determine which to associated with the following query

```
query GetSimulatedActivities(  
  $simulationDatasetId: Int!  
) {  
  simulated_activity(where: { simulation_dataset_id: {_eq: $simulationDatasetId } }) {  
    id  
    start_offset  
    duration  
    activity_type_name  
    attributes  
  }  
}
```

Linking an activity to a sequence is then a simple API call using the id from above

```
mutation LinkSimulatedActivityToSequence(  
  $seqId: String!  
  $simulationDatasetId: Int!  
  $simulatedActivityId: Int!  
) {  
  insert_sequence_to_simulated_activity_one(object: {  
    seq_id: $seqId  
    simulated_activity_id: $simulatedActivityId  
    simulation_dataset_id: $simulationDatasetId  
  }) {  
    seq_id  
  }  
}
```

Note: A simulated activity can only be associated with one sequence

Continue to [Retrieve SeqJson Serialization of Sequence](#)

Retrieve Sequence in SeqJson

Once a sequence is associated with activity instances, you can retrieve the SeqJson serialization of the sequence. This serialization is a simple concatenation of the commands associated with the activity instances in time-order. Any activities that do not have any associated commands (not yet expanded, no expansion associated, or some error during expansion) will not have any commands included in the resulting serialization. Additionally, any errors encountered during expansion will be included in the SeqJson as a command with the commands stem `$$ERROR$$`.

Retrieving the SeqJson serialization of a sequence is a simple API call

```
query GetSequenceSeqJson(  
  $seqId: String!  
  $simulationDatasetId: Int!  
) {  
  getSequenceSeqJson(  
    seqId: $seqId,  
    simulationDatasetId: $simulationDatasetId  
  ) {  
    id  
    metadata  
    steps {  
      type  
      stem  
      args  
      time {  
        tag  
        type  
      }  
      metadata  
    }  
  }  
}
```

Generate a standalone sequence in SeqJson

Some users will like to build a sequence seqJSON file outside of AERIE. Below is an example Sequence File that is compatible with AERIE. All you have to do is upload the contents of this file using the graphql command. The sequence seqJSON file that are generated are not saved into AERIE.

```
export default () =>
  Sequence.new({
    seqId: 'test_00001',
    metadata: {},
    commands: [
      BAKE_BREAD,
      A`2020-060T03:45:19`.PREHEAT_OVEN(100)
    ],
  });
```

GraphQL Query

```
query GetUserSequenceSeqJson($commandDictionaryID: Int!, $edslBody: String!) {
  getUserSequenceSeqJson(commandDictionaryID: $commandDictionaryID, edslBody: $edslBody) {
    id
    metadata
    steps {
      args
      metadata
      stem
      time {
        tag
        type
      }
      type
    }
  }
}
```

With the following `$edslBody` :

```
{
  "commandDictionaryID": 1,
  "edslBody": "export default () =>\n  Sequence.new({\n    seqId: \"test_00001\", \n    metadata: {}, \n    commands: [\n      BAKE_
BREAD, \n      A`2020-060T03:45:19`.PREHEAT_OVEN(100), \n    ], \n  });"
```

Results:

```
{
  "id": "test_00001",
  "metadata": {},
  "steps": [
    {
      "args": [],
      "metadata": {},
      "stem": "BAKE_BREAD",
      "time": {
        "type": "COMMAND_COMPLETE"
      },
      "type": "command"
    },
    {
      "args": [100],
      "metadata": {},
      "stem": "PREHEAT_OVEN",
      "time": {
        "tag": "2020-060T03:45:19.000",
        "type": "ABSOLUTE"
      },
      "type": "command"
    }
  ]
}
```

Valid Sequence JSON Format

We have discussed with the JSON Sequence Format team what will be the minimum needed to generate SCMFS for Europa Clipper.

In order to generate SCMFS, we need a Sequence JSON containing the mandatory fields below:

```
{
  "id": "<seq_id>",
  "metadata": {
    "onboard_name": "<onboard_name>",
    "onboard_path": "<onboard_path>"
  },
  "steps": [
    {
      "type": "command",
      "stem": "<command_stem>",
      "args": [] || [ args1, args2, etc.],
      "time": {
        "type": "COMMAND_COMPLETE"
      } || {
        "tag": "YYYY-DDDThh:mm:ss.sss",
        "type": "ABSOLUTE"
      } || {
        "tag": "hh:mm:ss.sss",
        "type": "COMMAND_RELATIVE"
      } || {
        "tag": "hh:mm:ss.sss",
        "type": "EPOCH_RELATIVE"
      }
    }
  ]
}
```

Real World Example with an Europa Clipper Command Dictionary

Here is an example expansion logic for a single Activity in a AERIE plan.

```
REAS_TX_ENABLE("DIS_CHIRP_DIS_RF", "DIS_CHIRP_DIS_RF"),
SEQ_VAR_CMD({ dynamic_arg1: "dynamic_2", opcode: 2 }).absoluteTiming(
  Temporal.Instant.from("2022-04-20T20:17:13Z")
),
DDM_DMP_EHA_PERIODIC({
  criteria_ops_cat1: "ALL",
  channel_id1: 1,
  criteria_ops_cat2: "AVS",
  channel_id2: 2,
  eha_evaluation_rate: "8 hz",
  execution_duration: 3,
  dp_priority: 62,
}).relativeTiming(Temporal.Duration.from({ minutes: 5 })),
AVS_SS_READ_CHECK_EMEM_NAND(10, 10, 10, false, 10),
DDM_DMP_EHA_PERIODIC({
  channel_id2: 2,
  channel_id1: 1,
  criteria_ops_cat2: "AVS",
  criteria_ops_cat1: "ALL",
  eha_evaluation_rate: "8 hz",
  execution_duration: 3,
  dp_priority: 62,
}).epochTiming(Temporal.Duration.from({ minutes: 15 })),
```

Below is the Sequence JSON file generate by AERIE:


```
{
  "id": "test0001",
  "metadata": {
    "onboard_name": "test.seq",
    "onboard_path": "/eng"
  },
  "steps": [
    {
      "type": "command",
      "stem": "REAS_TX_ENABLE",
      "args": ["DIS_CHIRP_DIS_RF", "DIS_CHIRP_DIS_RF"],
      "time": {
        "type": "COMMAND_COMPLETE"
      },
      "metadata": {}
    },
    {
      "type": "command",
      "stem": "SEQ_VAR_CMD",
      "args": ["dynamic_2", 2],
      "time": {
        "tag": "2022-110T20:17:13.000",
        "type": "ABSOLUTE"
      },
      "metadata": {}
    },
    {
      "type": "command",
      "stem": "DDM_DMP_EHA_PERIODIC",
      "args": ["ALL", 1, "AVS", 2, "8 hz", 3, 62],
      "time": {
        "tag": "00:05:00.000",
        "type": "COMMAND_RELATIVE"
      },
      "metadata": {}
    },
    {
      "type": "command",
      "stem": "AVS_SS_READ_CHECK_EMEM_NAND",
      "args": [10, 10, 10, 0, 10],
      "time": {
        "tag": "00:15:00.000",
        "type": "EPOCH_RELATIVE"
      },
      "metadata": {}
    }
  ]
}
```

User Guide

Overview

This document is a guide to how to make use of current Aerie capabilities. Aerie is a new software system to support the activity planning, sequencing and spacecraft analysis needs of missions. Aerie is being developed by the MPSA element of Multi-mission Ground System and Services (MGSS), a subsystem of AMMOS (Advanced Multi-mission Operations System). This guide will be updated as new features are added.

Aerie is a collection of loosely coupled services that support activity planning and sequencing needs of missions with modelling, simulation, scheduling and rule validation capabilities. Aerie will replace legacy MGSS tools including but not limited to APGEN, SEQGEN, MPS Editor, MPS Server, Sinc II / CTS and ULGEN. Aerie currently provides the following capabilities:

1. Merlin adaptation framework offering a subset of APGEN capabilities,
2. Merlin web GUI for activity planning,
3. Merlin command line interface for activity planning, and
4. Falcon smart sequence editor GUI.

Prerequisites

An adaptation is software developed with the Merlin framework libraries that models spacecraft behavior while performing a set of activities over a plan duration. Merlin adaptations can simulate a variety of States that are perturbed by executed activities, and governed by system models. Merlin plans describe a scheduled collection of activity instances with specified parameters. This user guide describes how to upload an existing adaptation .JAR file, how to create plans with that adaptation, and how to edit and simulate those plans. For users to complete these steps, they should be able to develop and compile an adaptation and have access to an Aerie installation. For details of how to create adaptations with Aerie refer to the [Mission Modeler Guide](#). For information on how to install Aerie services refer to the [Product Guide](#).

Merlin Activity Plans

Aerie provides an API to manage a database of plans. The database of plans may be queried for a list of all plans, and new plans may be added to the repository. Existing plans may be retrieved in full, replaced in full or in part, or deleted in full. The list of activities in a plan may be appended to (by creating a new activity) and retrieved in full. Individual activities in a plan may be retrieved in full, replaced in full or in part, and deleted in full. How to execute queries and mutations against the Aerie API is found in the [GraphQL software interface specification](#).

Operations on plans are validated to ensure consistency with the mission model-specific activity model with which they are associated. Stored plans shall contain activities whose parameter names and types are defined by the associated activity type.

Aerie GraphQL API

Purpose

This document describes the Aerie GraphQL API software interface provided by the latest [Aerie release](#).

Terminology and Notation

No special notation is used in this document.

Interface Overview

GraphQL is not a programming language capable of arbitrary computation, but is instead a language used to query application servers that have capabilities defined by the GraphQL specification. Clients use the GraphQL query language to make requests to a GraphQL service.

The three key GraphQL terms are;

- **Schema:** a type system which defines the data graph. The schema is the data sub-space over which queries can be defined.
- **Query:** a JSON-like read-only operation to retrieve data from a GraphQL service.
- **Mutation:** An explicit operation to effect server side data mutation.

A REST API architecture defines a particular URL endpoint for each "resources". In contrast, GraphQL's conceptual model is an entity graph. As a result, entities in GraphQL are not identified by URL endpoints and GraphQL is not a REST architecture API. Instead, a GraphQL server operates on a single endpoint, and all GraphQL requests for a given service are directed at this endpoint. Queries are constructed using the query language and then submitted as part of a HTTP request (either GET or POST).

The schema (graph) defines nodes and how they connect/relate to one another. A client composes a query specifying the fields to retrieve (or mutation for fields to create/update). A client develops their query/mutation with reference to the exposed GraphQL schema. As a result, a client develops custom queries/mutations targeted to its own use cases to fetch only the needed data from the API. In many cases this may reduce latency and increase performance by limiting client side data manipulation/filtering. For example, with a REST API there may be significant client side overhead and request latency when querying for an entire plan and then filtering the plan for the specific information of concern (and the work of adding filter fields as parameters to the end point). In contrast the GraphQL API allows the client to request only the fields of the plan data structure needed to satisfy the client's use case.

The Aerie GraphQL API is versioned with Aerie releases. However, a GraphQL based API gives greater flexibility to clients and Aerie when evolving the API. Adding fields and data to the schema does not affect existing queries. A client must specify the fields that make up their query. The additional fields in the graph simply play no part in the client's composed query. As a result, additions to the API do not require updates on the client side. However, clients do need to deal with schema changes when fields are removed or type definitions are evolved. Furthermore, GraphQL makes possible per-field auditing of the frequency and combinations with which certain fields are referenced by a client's queries and mutations. This provides Aerie developers with evidence of usage frequency which informs decision processes regarding deprecating or updating fields in the schema.

GraphQL Query Fundamentals

A round trip usage of the API consists of the following three steps:

1. Composing a request (query or mutation)
2. Submitting the request via POST
3. Receiving the result as JSON

POST Request

A standard GraphQL HTTP POST request should use the `application/json` content type, and include a JSON-encoded body of the following form:

```
{
  "query": "...",
  "operationName": "...",
  "variables": { "myVariable": "someValue" }
}
```

`operationName` and `variables` are optional fields. The `operationName` field is only required if multiple operations are present in the query.

Response

Regardless of the method by which the query and variables are sent, the response is returned in the body of the request in JSON format. A query's results may include some data and some errors, and those are returned in a JSON object of the form:

```
{
  "data": {},
  "errors": []
}
```

If there were no errors returned, the `"errors"` field is not present in the response. If no data is returned, the `"data"` field is only included if the error occurred during execution.

GraphQL Clients

Since a GraphQL API has more underlying structure than a REST API, there are a range of methods by which a client application may choose to interact with the API. A simple usage could use the `curl` command line tool, whereas a full featured web application may integrate a more powerful client library like [Apollo Client](#) or [Relay](#) which automatically handle query building, batching and caching.

Command Line

One may build and send a query or mutation via any means that enable an HTTP POST request to be made against the API. For example, this can be easily done using the command line tool [Graphqlurl](#).

GraphQL Console

The GraphQL API is described by a schema of the data graph. One can view the schema of the installed version of Aerie at http://<your_domain_here>:8080/console/api/api-explorer. The GraphQL console allows one to compose and test queries.

Browser Developer Console

Requests can also be tested using JavaScript from a web-browser. For example the following JavaScript can be used to make a simple query using the [fetch API](#):

```
const response = await fetch('http://<your_domain_here>:8080/v1/graphql', {
  body: JSON.stringify({ query: 'query { plan { id } }' }),
  headers: {
    Accept: 'application/json',
    Authorization: 'SSO_COOKIE_VALUE_HERE',
    'Content-Type': 'application/json',
  },
  method: 'POST',
});
const data = await response.json();
console.log(data);
```

Here is an example of a possible shape of the data logged to the console:

```
{
  "data": {
    "plans": [{ "id": 1 }, { "id": 2 }]
  }
}
```

This JavaScript can then be used as a hard-coded query within a client tool/script. For more complex and dynamic interactions with the Aerie API it is recommended to use a GraphQL client library.

Client Libraries

When developing a full featured application that requires integration with the Aerie API it is advisable that the tool make use of one of the many powerful GraphQL client libraries like [Apollo Client](#) or [Relay](#). These libraries provide an application functionality to manage both local and remote data, automatically handle batching, and caching.

In general, it will take more time to set up a GraphQL client. However, when building an Aerie integrated application, a client library offers significant time savings as the features of the application grow. One might choose to begin using HTTP requests as the client's API integration mechanism and later switch to a client library as the application becomes more complex.

GraphQL clients exist for the following programming languages:

- C# / .NET
- Clojurescript
- Elm
- Flutter
- Go
- Java / Android
- JavaScript
- Julia
- Kotlin
- Swift / Objective-C iOS
- Python
- R

A full description of these clients is found at <https://graphql.org/code/#graphql-clients>

Aerie GraphQL API

Schema

The schema is too large to include here and in Aerie's automatic documentation generation. The schema for your Aerie installation can be viewed at `http://<your_domain_here>:8080/console/api/api-explorer`.

Usage

Aerie employs the [Hasura](#) GraphQL engine to generate the Aerie GraphQL API. It is important to understand the significance and power of a data graph based API. A small primer of the GraphQL syntax can be found [here](#).

Queries & Subscriptions

- [Postgres: GraphQL Queries](#)
- [Postgres: GraphQL Subscriptions](#)
- [API Reference - Query / Subscription](#)

Mutations

- [Postgres: GraphQL Mutations](#)
- [API Reference - Mutation](#)

Examples

The following queries are examples of what Aerie refers to as "canonical queries" because they map to commonly understood use cases and data structures within mission subsystems. When writing a GraphQL query, refer to the schema for all valid fields that one can specify in a particular query.

Query for All Plans

```

query {
  plan {
    duration
    id
    model_id
    name
    start_time
  }
}

```

Query a Single Plan

```

query {
  plan_by_pk(id: 1) {
    duration
    id
    model_id
    name
    start_time
  }
}

```

Query for All Activity Instances for a Plan

You can either query `plan_by_pk` for all activity instances for a single plan, or query `plan` for all activity instances from *all* plans.

The following query returns all activity instances for a single plan:

```

query {
  plan_by_pk(id: 1) {
    activities {
      arguments
      id
      type
    }
  }
}

```

Query the Mission Model of a Plan

```

query {
  plan_by_pk(id: 1) {
    mission_model {
      id
      mission
      name
      owner
      version
    }
  }
}

```

Query for All Activity Types within a Mission Model of a Plan

Returns a list of activity types. For each activity type, the name, parameter schema, which parameters are required (must be defined).

```

query {
  plan_by_pk(id: 1) {
    mission_model {
      activity_types {
        computed_attributes_value_schema
        name
        parameters
        required_parameters
      }
    }
  }
}

```

Run Simulation

The following query starts a simulation, or returns information if the simulation has already started. The `simulationDatasetId` can be used to query for simulation results.

```

query {
  simulate(planId: 1) {
    reason
    simulationDatasetId
    status
  }
}

```

Query for All Simulated Activities in a Simulated Plan

```

query {
  plan_by_pk(id: 1) {
    simulations(limit: 1) {
      datasets(order_by: { id: desc }, limit: 1) {
        simulated_activities {
          activity_type_name
          attributes
          duration
          id
          parent_id
          simulation_dataset_id
          start_offset
        }
      }
    }
  }
  start_time
}

```

Query for All Resource Profiles in Simulated Plan

Profiles are simulated resources. The following query gets profiles for a given plan's latest simulation dataset (i.e. the latest resource simulation results):

```

query {
  plan_by_pk(id: 1) {
    duration
    simulations(limit: 1) {
      datasets(order_by: { id: desc }, limit: 1) {
        dataset {
          profiles {
            name
            profile_segments {
              dynamics
              start_offset
            }
          }
          type
        }
      }
    }
  }
  start_time
}

```

Query for All Simulated Activities and Resource Profiles in Simulated Plan

The following query just combines the previous two queries to get all activities and profiles in a simulated plan:


```

query {
  plan_by_pk(id: 1) {
    duration
    simulations(limit: 1) {
      datasets(order_by: { id: desc }, limit: 1) {
        simulated_activities {
          activity_type_name
          attributes
          duration
          id
          parent_id
          simulation_dataset_id
          start_offset
        }
        dataset {
          profiles {
            name
            profile_segments {
              dynamics
              start_offset
            }
          }
          type
        }
      }
    }
  }
  start_time
}

```

Query for All Resource Samples in Simulated Plan

```

query {
  resourceSamples(planId: 1) {
    resourceSamples
  }
}

```

Query for All Constraint Violations in Simulated Plan

```

query {
  constraintViolations(planId: 1) {
    constraintViolations
  }
}

```

Query for All Resource Types in a Mission Model

```

query {
  resourceTypes(missionModelId: 1) {
    name
    schema
  }
}

```

Create Plan

```

mutation {
  insert_plan_one(
    object: {
      duration: "432000 seconds 0 milliseconds"
      model_id: 1
      name: "My First Plan"
      start_time: "2020-001T00:00:00"
    }
  ) {
    id
    revision
  }
}

```

Create Simulation

Each plan must have a least one associated simulation to execute a simulation. To create a simulation for a plan you can use the following mutation:

```
mutation {
  insert_simulation_one(
    object: { arguments: {}, plan_id: 1, simulation_template_id: null }
  ) {
    id
  }
}
```

Create Activity Instances

```
mutation {
  insert_activity(
    objects: [
      {
        arguments: { peelDirection: "fromTip" }
        plan_id: 1
        start_offset: "1749:01:35.575"
        type: "PeelBanana"
      }
      {
        arguments: { peelDirection: "fromTip" }
        plan_id: 1
        start_offset: "1750:01:35.575"
        type: "PeelBanana"
      }
    ]
  ) {
    returning {
      id
      start_offset
    }
  }
}
```

Query for Activity Effective Arguments

This query returns a set of effective arguments given a set of required (and overridden) arguments.

```
query {
  getActivityEffectiveArguments(
    missionModelId: 1
    activityTypeName: "BakeBananaBread"
    activityArguments: { tbSugar: 1, glutenFree: false }
  ) {
    arguments
    errors
    success
  }
}
```

Resulting in:

```
{
  "data": {
    "getActivityEffectiveArguments": {
      "arguments": {
        "temperature": 350,
        "tbSugar": 1,
        "glutenFree": false
      },
      "success": true
    }
  }
}
```

When a required argument is not provided, the returned JSON will indicate which argument is missing. With `examples/banananation`'s `BakeBananaBread`, where only the `temperature` parameter has a default value:

```

query {
  getActivityEffectiveArguments(
    missionModelId: 1
    activityTypeName: "BakeBananaBread"
    activityArguments: {}
  ) {
    arguments
    errors
    success
  }
}

```

Results in:

```

{
  "data": {
    "getActivityEffectiveArguments": {
      "arguments": {
        "temperature": 350
      },
      "errors": {
        "tbSugar": {
          "schema": {
            "type": "int"
          },
          "message": "Required argument for activity \"BakeBananaBread\" not provided: \"tbSugar\" of type ValueSchema.INT"
        },
        "glutenFree": {
          "schema": {
            "type": "boolean"
          },
          "message": "Required argument for activity \"BakeBananaBread\" not provided: \"glutenFree\" of type ValueSchema.BOOLEAN"
        }
      },
      "success": false
    }
  }
}

```

Query for Mission Model Configuration Effective Arguments

The `getModelEffectiveArguments` returns the same structure as `getActivityEffectiveArguments`; a set of effective arguments given a set of required (and overridden) arguments. For example, `examples/config-without-defaults` 's has all required arguments:

```

query {
  getModelEffectiveArguments(missionModelId: 1, modelArguments: {}) {
    arguments
    errors
    success
  }
}

```

Results in:

```
{
  "data": {
    "getModelEffectiveArguments": {
      "arguments": {},
      "errors": {
        "a": {
          "schema": {
            "type": "int"
          },
          "message": "Required argument for configuration \"Configuration\" not provided: \"a\" of type ValueSchema.INT"
        },
        "b": {
          "schema": {
            "type": "real"
          },
          "message": "Required argument for configuration \"Configuration\" not provided: \"b\" of type ValueSchema.REAL"
        },
        "c": {
          "schema": {
            "type": "string"
          },
          "message": "Required argument for configuration \"Configuration\" not provided: \"c\" of type ValueSchema.STRING"
        }
      },
      "success": false
    }
  }
}
```

Simulation Results

After running a simulation, the results will be saved in the database, and queryable via the GraphQL API.

Querying for Simulation Results

There are two ways to query for simulation results: via the `simulate` action, and directly from the database.

Querying for simulation results via the `simulate` action

The `simulate` action always immediately returns a status. That status is one of the following:

- pending
- incomplete
- complete
- failed

The first time merlin receives a `simulate` request, it will add the request to its queue, and immediately return an "incomplete" status. The expectation is that the user will re-send the same request repeatedly, with an interval of at least 0.25 seconds, until one of the other two statuses is returned.

After the simulation completes, subsequent requests will receive either a `complete` or a `failed` status. If the simulation fails, the `reason` field will be populated with a message explaining what went wrong.

A message with a `complete` status will also populate the `results` field with information about the results of the simulation.

Querying for simulation results directly

In the near future, the recommended way to get simulation results will be directly from the database, rather than via merlin. Docs forthcoming as we iron out the details there.

There are several pieces of information included in the simulation results.

Contents of Simulation Results

Resource Profiles

For each resource, a profile allows the user to determine the value of that resource at any given time during the simulation.

There are currently two types of resource profiles, which each explain how to interpolate their values. A **discrete profile** is a list of values and times. The value of that resource at some time `t` would be the value of the last profile segment prior to `t`. A **real profile** allows smoother interpolation defined by a linear equation. Thus the value at time `t` is the value at the last profile segment prior to `t` (let's call it `t_prev`) plus the value of the linear equation evaluated at `t - t_prev`.

Resource Samples

See the `resourceSamples` GraphQL query for details on how to query for resource samples.

Simulated Activities

These are records of completed activity instances. They include not only the activities that were part of the plan, but also their child activities (see [decomposition](#)).

The **type** of a simulated activity is a string referring to the `ActivityType` of this activity instance.

The **arguments** is a map of strings to serialized values representing the arguments that were supplied to this activity instance. These arguments can come from the planner, or, in the case of child activities, from the parent activity.

The **start** instant represents the time at which this activity started. This time may have been set by the planner, or, in the case of child activities, it may have been determined by the logic in the parent's effect model.

The **duration** of an activity is the time that elapsed between the start of an activity and the end of that activity. Note that an activity is not considered to be terminated until all of its children have also terminated.

The **parentId** of an activity in the case of a child activity refers to the id of the parent activity. In the case of an activity created by a planner, this field will be **null**.

The **childIds** is a list of ids of the children of this activity. This list will be empty if this activity has not spawned any children.

The **directiveId** will be populated if this activity itself was placed by the planner. It will be empty if it was a child activity.

Note: **directiveId** and **parentId** are related concepts - if one is present, the other is absent, and vice versa.

The **computedAttributes** of an activity is a piece of information logged by the effect model upon termination of the activity.

The schema for the computed attributes can be found by querying the Activity Type, which is not included in simulation results. If the activity type does not define a computed attributes schema, the default value for computed attributes is **VOID**.

Unfinished Activities

Some activities may not terminate within the time limit of the simulation. The types and inputs to the activities will be listed here. Unfinished activities will not have computed attributes, since those are provided upon termination of an activity.

Events

Every change to the simulation is recorded in an event. Some of these events can be exported as part of simulation results, at the discretion of the mission model.

Planning UI

The Aerie planning web application provides a graphical user interface to create, view, update and delete mission model and plans.

Uploading Mission Models

Mission models can be uploaded to Aerie via the UI. To navigate to the [mission model page](#), click the **Mission Models** icon on the on the left navigation bar. Once an [mission model JAR](#) is prepared, it can be uploaded to the mission model service with a name, version, mission and owner. The name and version must match (in case and form) the name and version specified in the mission model.

For example, if the mission model is defined in code as `@MissionModel(name="Banananation", version="0.0.1")`, then the name field must be entered as **Banananation** and the version as **0.0.1**. Once the mission model is uploaded it will be listed in the table shown in Figure 1. Mission Models can be deleted from this table using the context menu by right clicking on the mission model.




Figure 1: Upload mission model, and view existing mission model.

Creating Plans

To navigate to the [plans page](#), click the **Plans** icon on the left navigation bar. Users can use the left panel to create new plans associated with any mission model. A **start** and **end** date has to be specified to create a plan. Existing plans are listed in the table on the right. Use right click on the table to reveal a drop down menu to delete and view plans.




Figure 2: Create plans, and view existing plans.

Base simulation arguments may be supplied at plan creation time. For more information please refer to [Aerie Simulation Configuration UI](#) documentation.

View and Edit Plans

Once a user clicks on an existing plan, they can view contents, add/remove activity instances, and edit activity instance parameters. The plan view is split into the following default panels:

- Schedule Visualization
- Simulation Visualization
- Activity Instances Table
- Side drawer containing:
 - Activity Dictionary
 - Activity Instance Details

In the default side drawer the activity dictionary is displayed. Once a type or instance is selected, users can view details such as metadata and parameters by moving the arrow keys down. Activities can be dragged into the timeline from the activity dictionary. Once instances are added they will appear in the Activity Instances table panel.

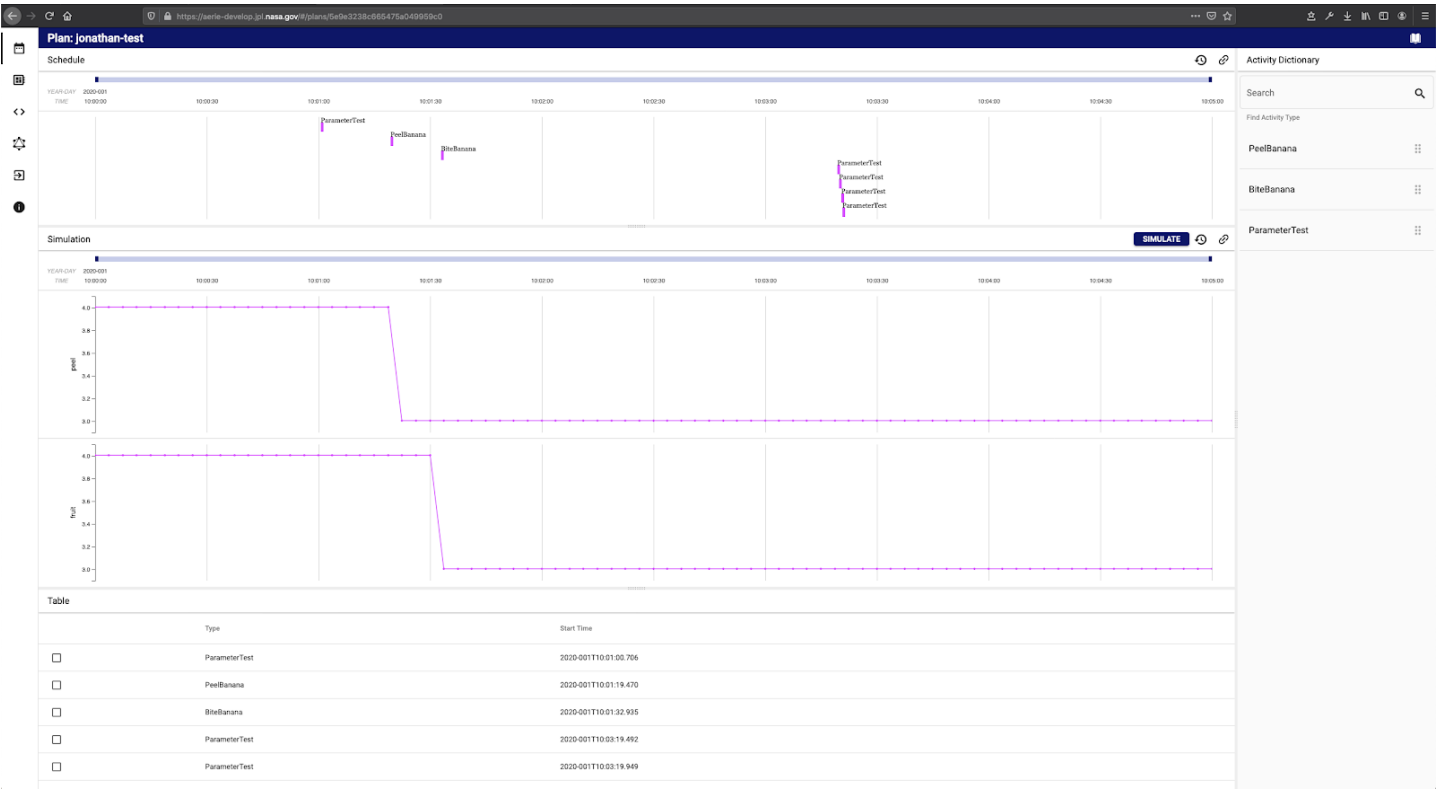


Figure 3: Default panels.

When a user clicks on an activity instance in the plan, the form to update activity parameters and start time will appear on the right drawer as shown in Figure 4. Users can use this form view to remove instances from the plan.

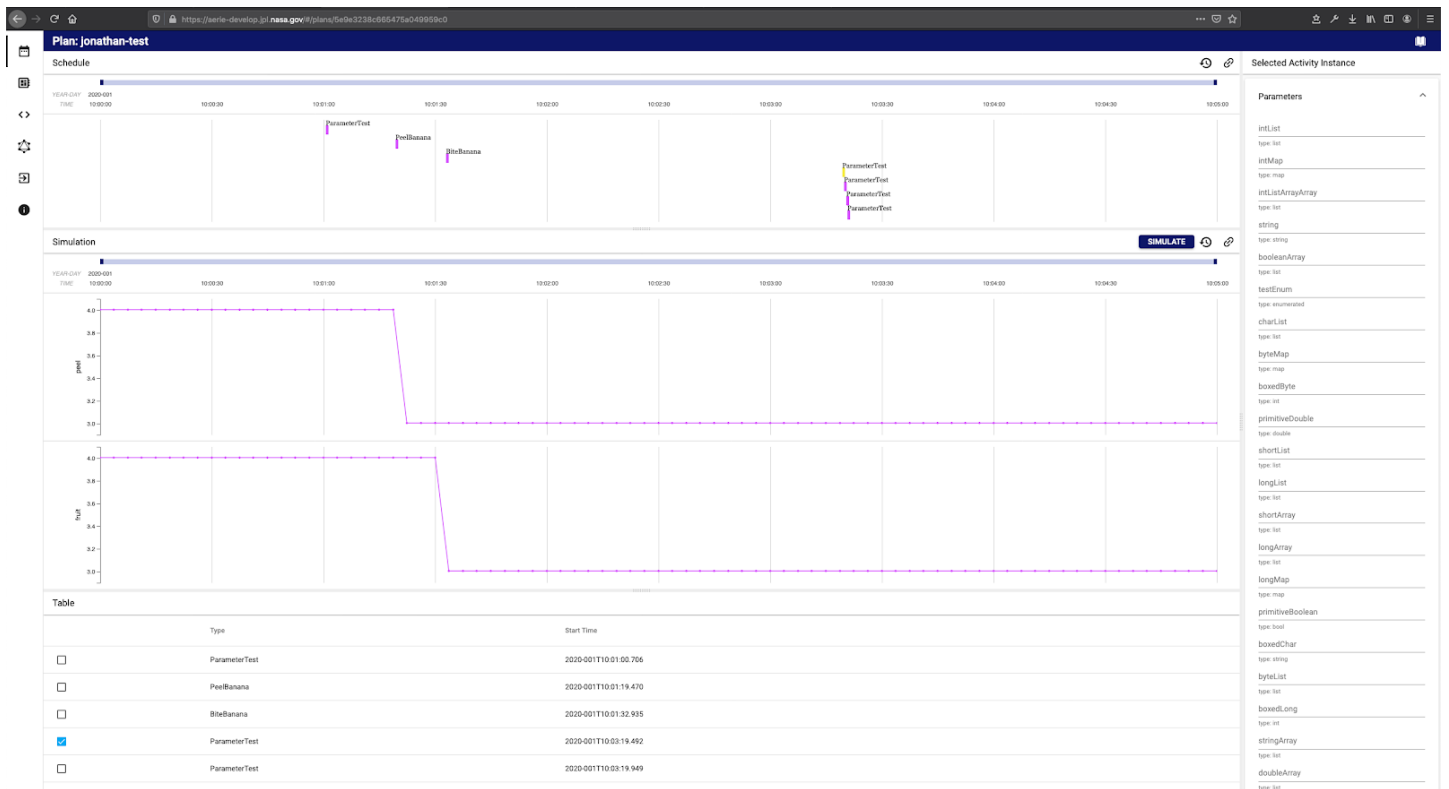


Figure 4: When an activity instance in plan is selected, its details will appear in the right drawer.

In the schedule and simulation elements, violations are shown as red regions in their respective bands as well as the corresponding time axis.

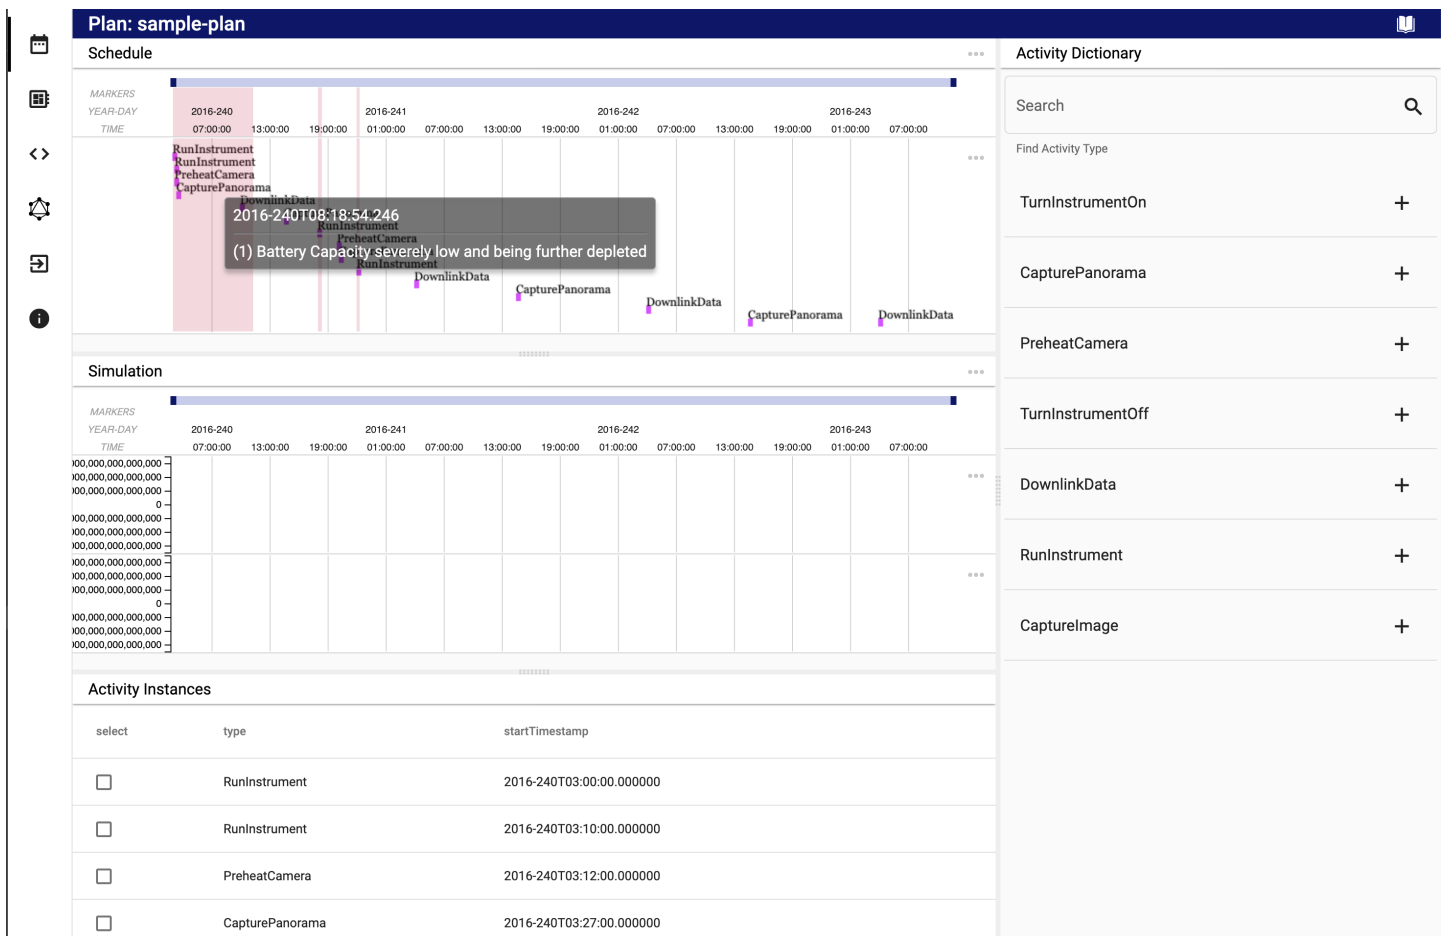



Figure 4.1: When violations occur, they will be represented as red areas on the timeline, with an accompanying hover state.

Horizontal guides may be added to any simulation or activity schedule band. The horizontal guides control UI may be accessed through each band's three dots more menu.



*Figure 4.2: Horizontal guides may be added to each activity schedule or simulation bands

Aerie UI provides a flexible arrangement where users can hide any of these panels by simply dragging dividers vertically. In Figure 5 this feature of the UI is illustrated.

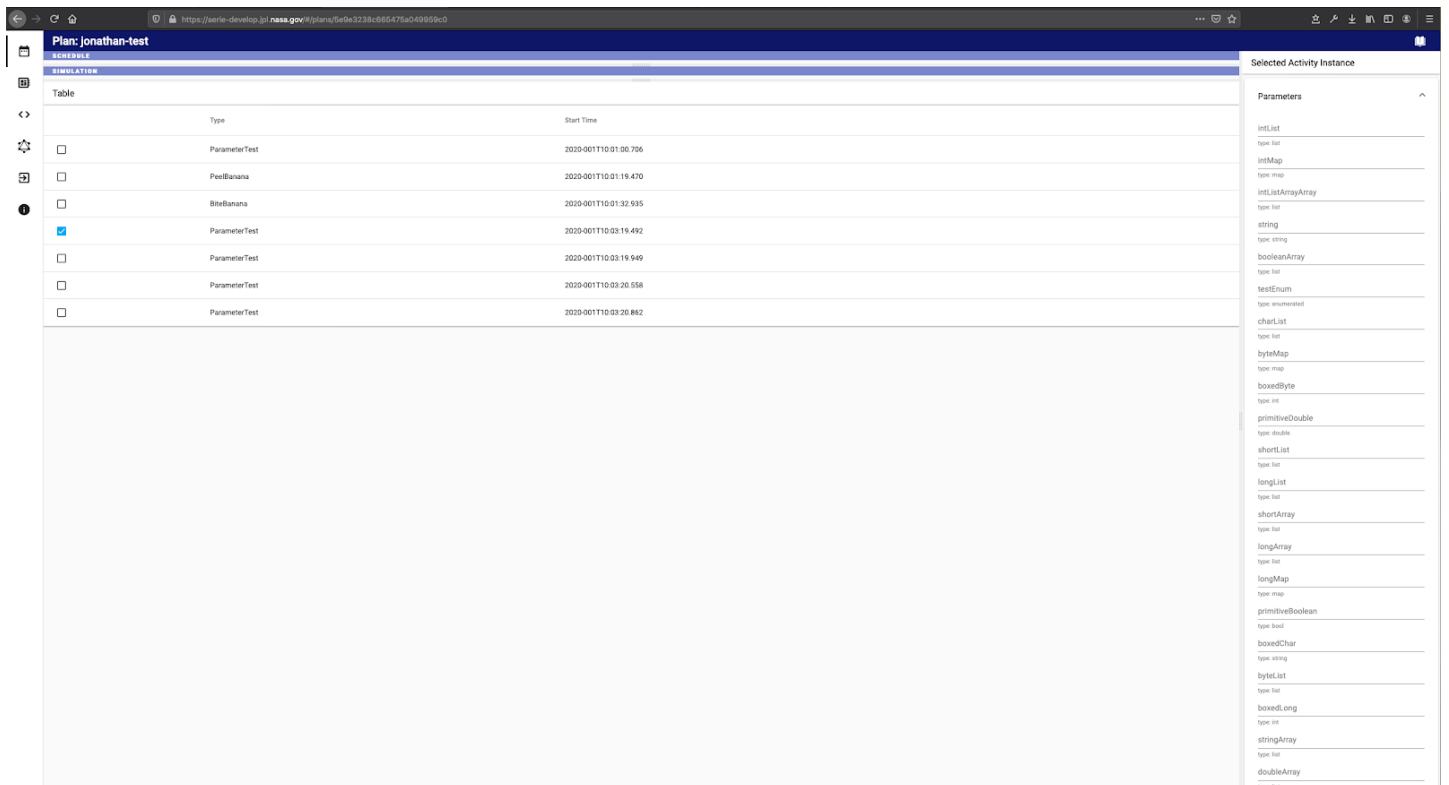


Figure 5: The bottom panels are dragged to the top edge completely leaving only one panels in view.

Note that all the default panels outlined here can be configured and changed based on the needs of a mission. You can read about how to do that in the [UI Configurability](#) documentation.

Simulation Configuration

The Aerie web application provides a graphical user interface to set and update simulation configuration arguments.

Using Web-App GUI

Navigate to the plans page. As seen in **Figure 1**, the plan creation form includes a "Simulation Configuration" section that accepts a user-provided file. Specifically, the user-provided file is expected to be a serialized (JSON) set of mission model arguments.




Figure 1: Plan creation and selection view.

Conceptually, these uploaded simulation configuration arguments are stored within a plan-specific simulation **template**. Within the GUI each plan may be associated with just one simulation template; individual simulations may set missing arguments or override existing arguments using the template's arguments as a base set of arguments. **Figure 2** shows the simulation configuration view where arguments for a specific simulation may be set/overridden.

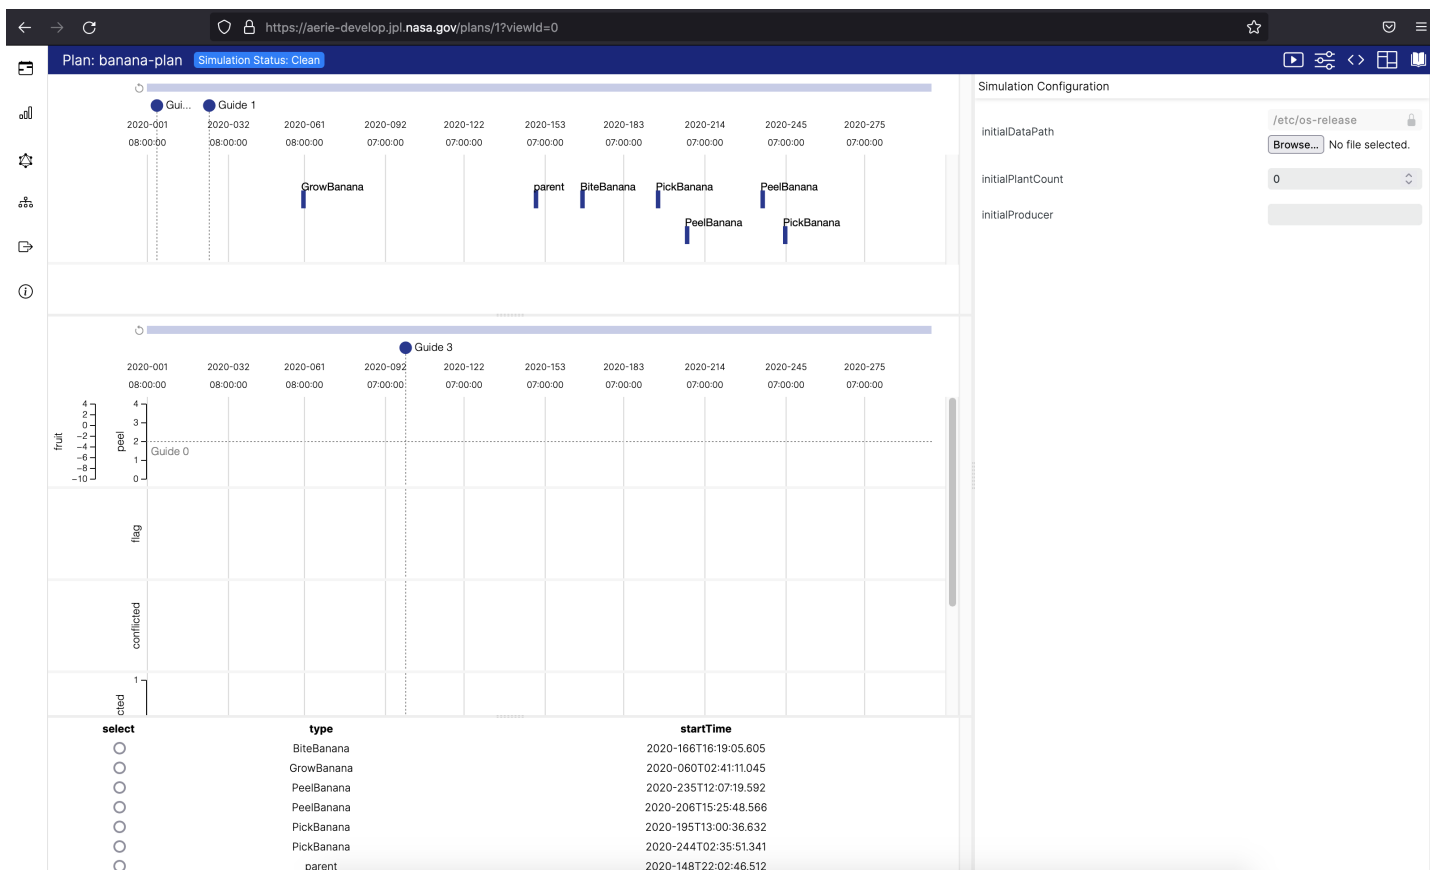


Figure 2: Simulation configuration view.

Formalized loosely, the sets of arguments at play here are:

- **T**, the template-defined arguments provided at plan creation time.
- **S**, the simulation-defined arguments provided prior to simulation invocation.
- **M**, the merged set of arguments derived from the union between **T** and **S**. Any duplicates will be resolved using arguments from **S**. This set is ultimately used in a simulation.

Using GraphQL API

Using the GraphQL API it is possible to define simulations from as many simulation templates as desired. Please refer to the [Aerie GraphQL API](#) [SIS](#) to get up-and-running with a GraphQL client capable of communicating with the Aerie API.

Create a Plan

Create a plan `test-plan`. This will trigger Merlin to create a new simulation that can be queried for from the GraphQL API:

```
query {
  simulation {
    id
  }
}
```

Which indicates that the one and only simulation `id` is:

```
{
  "data": {
    "simulation": [
      {
        "id": 1
      }
    ]
  }
}
```

This ID will be used to identify the newly created simulation.

Create a Simulation Template

A simulation template must be associated with a mission model. In this example the only existing mission model ID is `1`.

```
mutation {
  insert_simulation_template(objects: {
    model_id: 1,
    description: "first template",
    arguments: { initialPlantCount: 42, initialProducer: "template"}
  }) {
    returning {
      id
    }
  }
}
```

Resulting in:

```
{
  "data": {
    "insert_simulation_template": {
      "returning": [
        {
          "id": 1
        }
      ]
    }
  }
}
```

This ID will be used to identify the newly created simulation template.

Note that the arguments supplied here are not the full set of arguments required for simulation:

```
query {
  mission_model_parameters_by_pk(model_id: 1) {
    parameters
  }
}
```

Results in:

```
{
  "data": {
    "mission_model_parameters_by_pk": {
      "parameters": {
        "initialDataPath": {
          "order": 1,
          "schema": {
            "type": "path"
          }
        },
        "initialProducer": {
          "order": 2,
          "schema": {
            "type": "string"
          }
        },
        "initialPlantCount": {
          "order": 0,
          "schema": {
            "type": "int"
          }
        }
      }
    }
  }
}
```

As can be seen above, the template does not define a `initialDataPath` argument. In this case this must be provided by the simulation's argument set prior to simulation.

Associate Simulation with Simulation Template

Attach the simulation template to the current simulation and update the simulation's arguments.

```

mutation {
  update_simulation(
    _set: {
      simulation_template_id: 1,
      arguments: {
        initialPlantCount: 200,
        initialDataPath: "/etc/os-release"
      }
    },
    where: {id: {_eq: 1}}) {
    returning {
      id
    }
  }
}

```

The simulation with ID `1` is now associated with the template with ID `1`. In addition to being associated with a template, the simulation has defined a `initialPlantCount` argument override and a `initialDataPath` assignment.

Run the Simulation

```

query {
  simulate(planId: 1) {
    results
  }
}

```

The full results on this query are omitted for brevity but a sampling should look like:

```

"/plant": [
  {
    "x": 0,
    "y": 200
  },
  {
    "x": 31536000000000,
    "y": 200
  }
],
"/producer": [
  {
    "x": 0,
    "y": "template"
  },
  {
    "x": 31536000000000,
    "y": "template"
  }
]

```

Since an `initialProducer` argument was not provided in the simulation argument set the template's `initialProducer` value comes through here.

UI Views

Users can create custom planning views for different sub-systems (e.g. science, engineering, thermal, etc.), where only data (e.g. activities and resources) for those sub-systems are visualized. This is done through custom JSON files (or directly via the UI). The format of a UI View is the subject of this document.

See the [UI view JSON schema specification](#) for the complete set of view object properties and types. For more concrete examples see the [JSON view objects](#) included in the [default Aerie deployment](#).

View Schema

This is the main type interface for the planning UI view:

```
type ViewActivityTable = {
  columnDefs: ColumnDef[];
  id: number;
};

type ViewIFrame = {
  id: number;
  src: string;
  title: string;
};

type ViewDefinition = {
  plan: {
    activityTables: ViewActivityTable[];
    iFrames: ViewIFrame[];
    layout: Grid;
    timelines: Timeline[];
  };
};
```

For example, here is a JSON object that implements an empty `ViewDefinition` interface:

```
{
  "plan": {
    "activityTables": []
    "iFrames": [],
    "layout": {},
    "timelines": []
  }
}
```

Layout (Grid)

A planning UI view consists of a layout which describes the different visible components and how they are arranged in re-sizeable rows and columns. The layout follows the following `Grid` type definitions:

```

type GridComponent = {
  activityTableId?: number;
  componentName: string;
  gridName?: string;
  iFrameId?: number;
  id: number;
  props?: any;
  timelineId?: number;
  type: 'component';
};

type GridColumns = {
  columnSizes: string;
  columns: Grid[];
  gridName?: string;
  id: number;
  type: 'columns';
};

type GridGutter = {
  gridName?: string;
  id: number;
  track: number;
  type: 'gutter';
};

type GridRows = {
  gridName?: string;
  id: number;
  rowSizes: string;
  rows: Grid[];
  type: 'rows';
};

type Grid = GridColumns | GridComponent | GridGutter | GridRows;

```

For example, here is a grid layout definition in JSON:

```

{
  "columnSizes": "1fr 3px 2fr 3px 1fr",
  "columns": [
    { "componentName": "ActivityFormPanel", "id": 1, "type": "component" },
    { "id": 2, "track": 1, "type": "gutter" },
    {
      "id": 3,
      "rowSizes": "70% 3px 1fr",
      "rows": [
        {
          "componentName": "TimelinePanel",
          "id": 4,
          "timelineId": 0,
          "type": "component"
        },
        { "id": 5, "track": 1, "type": "gutter" },
        {
          "activityTableId": 0,
          "componentName": "ActivityTablePanel",
          "id": 6,
          "type": "component"
        }
      ]
    },
    { "id": 7, "track": 3, "type": "gutter" },
    { "componentName": "ActivityTypesPanel", "id": 8, "type": "component" }
  ],
  "gridName": "Activities",
  "id": 0,
  "type": "columns"
}

```

Timeline

The `timelines` section allows you to specify a list of `timeline` visualizations which display time-ordered data (i.e. activities or resources). Here is the interface of a timeline:


```
interface Timeline {
  id: number;
  marginLeft: number;
  marginRight: number;
  rows: Row[];
  verticalGuides: VerticalGuide[];
}
```

To visualize data in a timeline you need to add row objects to the `rows` array. A row is a layered visualization of time-ordered data. Each layer of a row is specified as an object of the `layers` array. The interfaces for a `Row` and `Layer` are as follows:

```
interface Row {
  autoAdjustHeight: boolean;
  height: number;
  horizontalGuides: HorizontalGuide[];
  id: number;
  layers: Layer[];
  yAxes: Axis[];
}

interface Layer {
  chartType: 'activity' | 'line' | 'x-range';
  filter: {
    activity?: ActivityLayerFilter;
    resource?: ResourceLayerFilter;
  };
  id: number;
  yAxisId: number | null;
}
```

Here is a JSON object that creates a single row with one activity layer. Notice the `filter` property, which is a [JavaScript Regular Expression](#) that specifies we only want to see `activity` of `type` `.*`. This is a regex for giving all activity types.

```
{
  "autoAdjustHeight": true,
  "height": 200,
  "horizontalGuides": [],
  "id": 0,
  "layers": [
    {
      "activityColor": "#283593",
      "activityHeight": 20,
      "chartType": "activity",
      "filter": { "activity": { "type": ".*" } },
      "id": 0,
      "yAxisId": null
    }
  ],
  "yAxes": []
}
```

For data that has y-values (for example resource data), you can specify a y-axis and link a layer to it by ID. Here are the interfaces for `Axis` and `Label`:

```
interface Axis {
  color: string;
  id: number;
  label: Label;
  scaleDomain: (number | null)[];
  tickCount: number | null;
}

interface Label {
  align?: CanvasTextAlign;
  baseline?: CanvasTextBaseline;
  color?: string;
  fontFace?: string;
  fontSize?: number;
  hidden?: boolean;
  text: string;
}
```

Y-axes are specified in the row separately from layers so we can specify multi-way relationships between axes and layers. For example you could have many layers corresponding to a single row axis.

Here is the JSON for creating a row with two overlaid `resource` layers. The first layer shows only resources with the name `peel`, and uses the y-

axis with ID `1`. The second layer shows only resources with the name `fruit`, and uses the y-axis with the ID `2`.

```
{
  "autoAdjustHeight": false,
  "height": 100,
  "horizontalGuides": [],
  "id": 1,
  "layers": [
    {
      "chartType": "line",
      "filter": { "resource": { "name": "peel" } },
      "id": 1,
      "lineColor": "#283593",
      "lineWidth": 1,
      "pointRadius": 2,
      "yAxisId": 1
    },
    {
      "chartType": "line",
      "filter": { "resource": { "name": "fruit" } },
      "id": 2,
      "lineColor": "#ffcd69",
      "lineWidth": 1,
      "pointRadius": 2,
      "yAxisId": 2
    }
  ],
  "yAxes": [
    {
      "color": "#000000",
      "id": 1,
      "label": { "text": "peel" },
      "scaleDomain": [0, 4],
      "tickCount": 5
    },
    {
      "color": "#000000",
      "id": 2,
      "label": { "text": "fruit" },
      "scaleDomain": [-10, 4],
      "tickCount": 5
    }
  ]
}
```

Activity Tables

Here is an example `activityTables` view JSON definition. The `columnDefs` and `columnStates` follow the schemas of the `ag-grid` `ColumnDef` and `ColumnState` respectively.

```
{
  "activityTables": [
    {
      "columnDefs": [
        { "field": "id", "filter": "agTextColumnFilter", "headerName": "ID", "sortable": true, "resizable": true },
        { "field": "type", "filter": "agTextColumnFilter", "headerName": "Type", "sortable": true, "resizable": true },
        { "field": "start_time", "filter": "agTextColumnFilter", "headerName": "Start Time", "sortable": true, "resizable": true },
        { "field": "duration", "filter": "agTextColumnFilter", "headerName": "Duration", "sortable": true, "resizable": true }
      ],
      "columnStates": [],
      "id": 0
    }
  ]
}
```

To use the `ActivityTablePanel` you need to add an `ActivityTablePanel` component to the grid layout and connect it via the `activityTableId`. For example:

```
{ "activityTableId": 0, "componentName": "ActivityTablePanel", "id": 2, "type": "component" }
```

Notice how we connect the grid component `activityTableId` with the `id` of the definition in the `activityTables` array.

IFrames

An `IFrame` component allows you to embed another application in the Aerie UI. Here is an example `iFrames` view JSON definition:

```
{
  "iFrames": [
    {
      "id": 0,
      "src": "https://iss-sim.spacex.com/",
      "title": "ISS-Docking-Simulator"
    }
  ]
}
```

To use the `IFrame` you need to add an `IFrame` component to the grid layout and connect it via the `iFrameId`. For example:

```
{ "componentName": "IFrame", "iFrameId": 0, "id": 1, "type": "component" }
```

Notice how we connect the grid component `iFrameId` with the `id` of the definition in the `iFrames` array.

GraphQL Queries

The following queries can be used to programmatically operate on UI views. For more information on how to use these queries please see the [Aerie GraphQL API documentation](#).

Create Single View

```
mutation {
  insert_view_one(
    object: {
      definition: {
        plan: { activityTables: [], iFrames: [], layout: {}, timelines: [] }
      }
      name: "My First View"
      owner: "system"
    }
  ) {
    id
  }
}
```

Create Multiple Views

```
mutation {
  insert_view(
    objects: [
      {
        definition: {
          plan: { activityTables: [], iFrames: [], layout: {}, timelines: [] }
        }
        name: "First View"
        owner: "system"
      },
      {
        definition: {
          plan: { activityTables: [], iFrames: [], layout: {}, timelines: [] }
        }
        name: "Second View"
        owner: "system"
      }
    ]
  ) {
    returning {
      id
    }
  }
}
```

Get All Views

```
query {
  view {
    created_at
    definition
    id
    name
    owner
    updated_at
  }
}
```

Get Single View by ID

```
query {
  view_by_pk(id: 1) {
    created_at
    definition
    id
    name
    owner
    updated_at
  }
}
```

Delete Single View by ID

```
mutation {
  delete_view_by_pk(id: 1) {
    id
  }
}
```

Delete All Views in the Database

```
mutation {
  delete_view(where: {}) {
    affected_rows
  }
}
```

Update Definition and Name of a Single View by ID

```
mutation {
  update_view_by_pk(
    pk_columns: { id: 1 }
    _set: { definition: { plan: {} }, name: "New Name" }
  ) {
    id
  }
}
```

Activity Directive Metadata

This document describes how to add mission-specific metadata (aka annotations) to activity directives. There are 6 different types of metadata which are summarized in the table below.

Type	Schema	Display
string	{ "type": "string" }	Single line text field
long_string	{ "type": "long_string" }	Textarea
boolean	{ "type": "boolean" }	Checkbox
number	{ "type": "number" }	Single line numeric field
enum	{ "enumerates": [], "type": "enum" }	Single select dropdown
string	{ "enumerates": [], "type": "enum_multiselect" }	Multi select dropdown

Adding Metadata Schemas

To add mission-specific metadata schemas you need to use the GraphQL API. Here is a mutation that adds an example of each type.

Mutation

```
mutation CreateActivityDirectiveMetadataSchemas($schemas: [activity_directive_metadata_schema_insert_input!]) {
  insert_activity_directive_metadata_schema(objects: $schemas) {
    affected_rows
    returning {
      created_at
      key
      schema
      updated_at
    }
  }
}
```

Query Variable

```
{
  "schemas": [
    { "key": "STRING_EXAMPLE", "schema": { "type": "string" } },
    { "key": "LONG_STRING_EXAMPLE", "schema": { "type": "long_string" } },
    { "key": "BOOLEAN_EXAMPLE", "schema": { "type": "boolean" } },
    { "key": "NUMBER_EXAMPLE", "schema": { "type": "number" } },
    { "key": "ENUM_EXAMPLE", "schema": { "enumerates": ["A", "B", "C"], "type": "enum" } },
    { "key": "ENUM_MULTISELECT_EXAMPLE", "schema": { "enumerates": ["D", "E", "F", "G"], "type": "enum_multiselect" } }
  ]
}
```

Product Guide

- [Product Installation](#)
- [Administration](#)
- [System Requirements](#)
- [Product Support](#)

Product Installation

Installation instructions are located in the Aerie repository [deployment documentation](#). There you can also find an example [docker-compose.yml](#) file.

Docker Images

Aerie Docker images are available at [GitHub packages](#).

Known Issues

1. Using a custom base URL path for Aerie UI
 - Follow the instructions linked at [Custom Base Path](#) to deploy Aerie UI at a custom base path.
2. Annotations processing being used for Activity Mapping
 - When using the IntelliJ IDE, upon a source file change, only the affected source files will be recompiled. This causes conflicts with the annotations processing being used for Activity Mapping. For now manually rebuilding every time is the solution.

Administration

This product is using Docker containers to run the application. The Docker containers are internally bridged (connected) to run the application. Containers can be restarted in case of any issues using Docker CLI. See the [TCP Port Requirements](#) for which containers should be exposed publicly/outside the Docker network.

Configuration

Each Aerie service is configured with environment variables. A description of those variables is found in the [Environment Variable Documentation](#)

Of note, the aerie-merlin, aerie-scheduler, and aerie_merlin_worker(s) can be provided additional JVM arguments as environment variables. For example one may choose to configure the JVM allocated heap size. One must provide any desired JVM flags to the `JAVA_OPTS` environment variable for the container being configured.

Network Communications

The Aerie deployment configures the port numbers for each container via docker-compose. The port numbers must match those declared within the services' config.json. In a large majority of Aerie deployments no change to these port numbers will be needed, nor should one be made. The only port number that might be desired to change is the Aerie-UI port (80). In this case the number to change is the first port number of the pair [XXXX:XXXX]. The second number represents the port number within the container itself.

Administration Procedures

Aerie employs an orchestrated containerized architecture. Each of the software components are independently packaged and run in a container.

- aerie-commanding: Provides activity to command expansion authoring and processing.
- aerie-gateway: Main API gateway for Aerie
- aerie-merlin: Handles all the logic and functionality for activity planning.
- aerie-merlin-worker: Runs simulation jobs requested of the aerie-merlin service.
- aerie-scheduler: Provides automated activity plan scheduling.
- aerie-ui: Hosts the web application and communicates with Aerie via the GraphQL Apollo Server.

- postgres: Holds the data for the Merlin server container. This container is optional and included in many of the aerie docker compose files for convenience. Long term deployments will want to integrate a Postgres database not running within a container.

- hasura: Serves the Aerie GraphQL API.

Aerie database containers are isolated and connect only with service containers internal to the application (Merlin server, UI server, Aerie gateway server). Aerie databases are not (and should not be made to be) accessible from outside the Aerie application.

Network File System Deployment

You can install and use a Network File System with Aerie. Please see the [deployment documentation](#) for complete instructions.

System Requirements

Software Requirements

Name	Version
Docker	19.X
PostgreSQL	14.X

Supported Browsers

Name	Version
Chrome	Latest
Firefox	Latest

Hardware Requirements

Hardware	Details
CPU	2 Gigahertz (GHZ) or above
RAM	8 GB at minimum
Storage	15 GB (the system should have access to additional storage as system databases grow according to mission operations)
Display resolution	2560-BY-1600, recommended
Internet connection	High-Speed connection, at least 10MBPS

AWS EC2 Instance Type

The workload that can be submitted to Aerie is highly dependent on the computational complexity of the mission model being simulated. An [m4.large](#) or greater EC2 instance will satisfy generic usage of Aerie with simple mission model. For missions that develop more complex mission models, operations such as performing simulation will benefit from the increased CPU of a [c4.xlarge](#) or greater instance.

TCP Port Requirements

Service	Default Port	Public
Aerie Commanding	27184	No
Aerie Gateway	9000	Yes
Aerie Merlin	27183	No
Aerie Merlin Worker	27187	No

Service	Default Port	Public
Aerie Scheduler	27185	No
Aerie UI	80	Yes
Hasura	8080	Yes
Postgres	5432	No

Product Support

Defect Reporting Procedure

All defect reports should go to aerie_support@jpl.nasa.gov.

Points of Contact

- Deployment and Integration: Kenneally, Patrick W, Development Lead
- Administration: Kenneally, Patrick W, Development Lead
- General Help: Alper Ramaswamy, Emine Basak, Product Lead